


VIRTUAL MEMORY CHAPTER # 10

→ Benefits in program execution that is only partially in memory.

- ① A program would no longer be constrained by the amount of physical mem that is available.
- ② Users would be able to write for an extreme large virtual address space
- ③ More programs run at same time
- ④ Incr CPU utilization & throughput
- ⑤ Less I/O would be needed to load or swap portions of program int mem.

→ The virtual address space of a process refers to a logical (or virtual) view of how a process is stored in memory.




pages can be shared during process creation through system call.

→ like indep # of

→ Virtual memory also used in storing memory b/w process

virtual address space that include holes are known as sparse address space. These holes can be filled when needed by a stack or heap.



MEMENTO PAGING

→ Storing mein humein pura program memory mein load krne ki need nhi.

→ The best strategy is to load pages only when they are needed.

→ Pages that are never accessed are never loaded in main memory.

→ While a process is executing some pages may be in main memory & some in secondary memory or backing store.

→ Valid / Invalid bits are used to distinguish b/w the pages which are in memory or in secondary.
Valid → Page is in main memory.

Invalid → Either the page is not in address space of process or it is valid but currently in ~~secon~~ backing store (secondary storage).

→ Access to page marked INVALID causes page fault.

→ PAGE FAULT HANDLING:

→ First check a valid or invalid memory access.

↳ if Invalid, then terminate the process

↳ else if valid then we page it in

→ Find a free frame (from a free frame list)

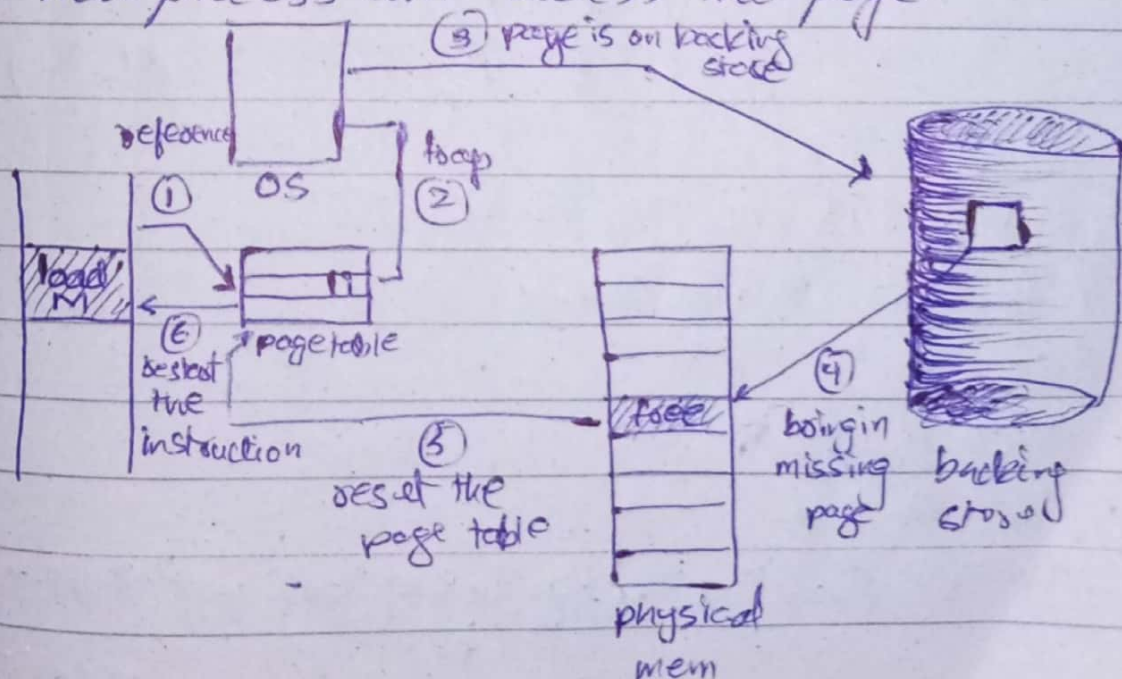
→ We now read the desired page from the backing store

→ The page is now allocated to the free frame.

→ The pagetable is then updated, to indicate that the page is now available in memory it is set to VALID.

→ We restart the instruction that was interrupted due to page fault.

→ Now process can access that page.



Thrashing: a phenomenon where the performance of a system significantly degrades due to excessive paging or swapping activity. It occurs when OS spends majority of its time moving pages b/w RAM & Secondary storage rather than executing actual process.

- pure demand pages: Never bring a page into memory until it is required.

- ↳ We can start executing a process with no page

- ↳ When OS sets the IP on the 1st instruction of process that page would not be available in main memory.

- ↳ This will result in page fault.

- ↳ Then page is brought into memory & continues its execution.

- ↳ Page faulting occurs unless every page is loaded in main memory.

- Hardware Support:

- ① Page Table

- ② Secondary Memory → This memory holds those pages that are not present in main memory. A part of SM i.e. swap space is specially allocated for this purpose.

- Restarting Instructions on Page Faults

1. Fetch and decode (ADD)

2. Fetch A

3. Fetch B

4. Add A & B

5. Store the Sum in C

→ After 5 steps per page fault occurs. For example if memory main load nhi hai to hame step 1 se restart krna parega.

- Reasons of thrashing:
- ① High degree of multiprogramming
 - ② Large working sets
 - ③ Insufficient physical mem

• Another Scenario: For example a ^{instruction} program is executing & during its execution it modifies or overwrites several data. After changing multiple data a page fault occurs. Now when we restart the instructions how can we get the initial values prior to modification.

Solution #01: The OS checks the source & destination operands whether they are loaded or not, if not a page fault occurs before an modification.

Solution #02: Use temporary registers to hold values of modified data. If there is a page fault the initial values are written back in memory.

Free Frame List: ^{needs} → [7] → [8] → [20] → ... [98]

→ To resolve page fault most OS uses a free-frame list.

→ Each frame is 'zeroed-out' before being allocated, thus erasing their previous content. This is called 'zero fill on demand'.

• Pre-Paging: OS guesses in advance which pages will be needed and pre-loads them into memory.

↳ If the guess of OS is wrong → page fault

↳ Errors may occur in removing useful pages

↳ Difficult to get right guess due to branch code.

"Performance on Demand Paging"

→ let m_a = memory access time

p = probability of page fault.

$$\text{effective access time} = (1-p)m_a + p \times \text{page fault time}$$

↳ means page is in memory
- may be in TLB
- or TLB miss
- Then page table entries loaded from main mem

↳ page fault probability

TLB miss
page not in mem
read from
swap file

→ With an avg page fault service time of 8ms and a memory access time of 200ns,

$$= (1-p) \times 200 + p \times 8\text{ms}$$

$$= (1-p) \times 200 + p \times 8,000,000$$

$$= 200 + 7,999,800 \times p \text{ ns}$$

effective access time \propto page fault.

→ if want performance degradation to be less than 10%,

$$220 > 200 + 7,999,800 \times p$$

$$20 > 7,999,800 \times p$$

$$p < 0.0000025$$

→ I/O to swap space is usually faster than filesystems.

→ The use of swap space also affects the performance.
Option #01: Copy entire process to swap space at program's startup than perform demand paging from swap space.


Option #02: (Better) Demand page from file system initially but to write the pages to swap space as they are replaced.

→ The demand paging of binary executable file is done from file system. As these files are not modified so they need to not be replaced again in secondary memory, they are simply replaced by another frame. And when needed again they are accessed from filesystem.

"COPY ON WRITE"

→ Copy-on-write allows the parent & child to share the same address space.

→ These shared pages are marked as copy-on-write, means make a copy of the page when ^{ever} either of the process ~~writes~~ tries to write on it.



→ Only pages that can be modified should be masked as COW.

→ Read-only or executable code can be shared by two process.

`vfork()`; Virtual memory fork.

→ During execution of child parent is suspended.

→ Child uses address space of parent

→ Does not use copy on write.

→ Any changes made by child will be visible to parent once it continues.

→ `vfork()` is intended to use when child calls `exec()` immediately after creation.

→ `vfork()` extremely efficient because of no copying of pages.

"PAGE REPLACEMENT"

→ Frame Allocation problem refers to the challenge of allocating limited physical memory frames among multiple process or pages in a virtual memory system.

↳ It involves deciding how many physical memory frames to allocate to each process or page & which frames should be allocated for efficient usage.

→ The page replacement problem refers to selection of a page to be evicted from main memory when a page fault occurs. This is done to add a new page to from secondary storage to main mem.

→ Following steps are taken when a page fault occurs. ~~Modified~~

① Find the location of the desired page on Secondary storage

② Find a free frame:

↳ if there is a free frame use it

↳ else if there is no free frame, use a page-replacement algo to select a victim frame.

↳ Write victim frame to secondary storage (if necessary); change the page & frame table

- ③ Read the desired page into the newly freed frame; change the page & frame table
- ④ Repeat the process.

→ The swapping of pages increase the effective access time.

Dirty Page: A dirty page in a memory is the one which has been modified (written to) since it was last loaded.

↳ So if a dirty page is replace it is needed to be written back to secondary storage.

Clean Page: A page which isn't modified since it was loaded. like executable or read-only file.

↳ These pages aren't required to be written back, they are simply discarded.

→ A modify bit (or dirty bit) is used for this purpose of this bit set means dirty page if clear means clean.

→ These scheme significantly reduces the time required to service a page fault as it reduces I/O to a secondary storage by one-half for a clean page.

① FIFO Page Replacement :

→ The first page brought in is replaced.

7	0	1	2	0	3	0	4	2	3	0	3	2	1	2	0	1	7
7	7	7	2	2	2	2	4	4	4	0		0	0				7
	0	0	0	0	3	3	3	2	2	2		1	1				1
		1	1	1	1	0	0	0	3	3		3	2				2
PF	PF	PF	PF		PF	PF	PF	PF	PF			PF	PF				PF

7 is replaced

1 2 3 4 1 2 5 1 2 3 4 5

→ For the above reference string if we perform FIFO for 3-frames we get
hit = 3 & page-faults = 9

But if we do with 4-frames we get,
hit = 2 & page-faults = 10

→ The page faults should have decreased by increasing no. of frames. This is called Belady's anomaly.

$$\Rightarrow \text{Hit ratio} = \frac{\text{No. of hits}}{\text{total no. in reference string}} = \frac{9}{18.6} = \frac{9}{10}$$

$$\Rightarrow \text{Page fault ratio} = \frac{\text{No. of faults}}{\text{No. of reference string}}$$

② Optimal Page Replacement :

"Replace the page that will not be used for the longest period of time"

→ This algo guarantees the lowest possible page fault rate for a fixed no. of frames.

now 1 is replace

7	0	1	2	0	3	0	4	2	3	0	3	2	1	2	0	1	7	0
7	7	7	2		2		2		2		2		2				7	
	0	0	0		0		0		0		0		0				0	
		1	1		3		3		3		1		1				1	
*	*	*	*	*	*	*	*	*	*	*	*	*	*	*	*	*	*	*

7 replace hua q k wo sabse last mein aya 7, 0, 1 mein se

→ This algo cannot be implemented as it requires future knowledge.

→ Hit ratio = $\frac{11}{20}$, Page fault ratio = $\frac{9}{20}$

→ If we reverse this reference string we will get same no. of page faults

③ LRU Page Replacement :

"Replace the page that has not been used for longest period of time."

→ Least Recently Used (LRU) algo.

→ It associates a timestamp with each page when it was last used.

→ Ismen hum peche dekhenge k kon sabse last use hua hai.

7	0	1	2	0	3	0	4	2	3	0	3	2	1	2	0	1	7	0	1
7	7	7	2		2		4	4	4	0				1		1		1	
	0	0	0		0		0	0	3	3				3		0		0	
		1	1		3		3	2	2	2				2		2		7	
×	×	×	×		×		×	×	×	×				×		×		×	

→ Hit ratio = $\frac{8}{20}$, Page fault = $\frac{12}{20}$

→ if we reverse the string we get same no. of page faults

→ Optimal & LRU are called Stack Algorithms, they never exhibit Belady's anomaly.

→ Timestamp, doubly linked list & stack.

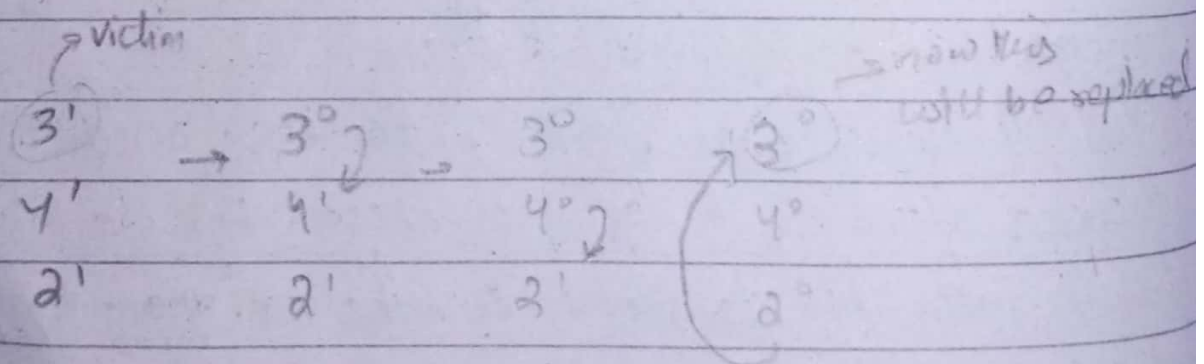
→ Stack Implementation: Jo koi page hoga recently referenced wo stack k top par hoga. Agr wo middle of stack mein hai to uska se uska kr. upper top par rak henge. → To expensive as we need to change all pointers.

→ Timestamp also costly as OS needs to look at all pages & update their timestamps.

④ Second Chance Algorithm:

- A modified form of FIFO page algo.
- Har page k saath ek reference bit hogi jo initially zero hogi.
- If the page is referenced again we set the bit.
- The page which has a zero reference bit, is ^{has to replace} replaced.
- Agr saari hi reference bit ~~zero~~ 1 hai to sabki
- Agr js page ko replace krna hai uski ref bit 1 hai to usko hum zero krdenge or jo next replace ^{hone} ~~krne~~ wala hoga uski check krenge. Agr wo 0 hai to replace agr 1 ho to next wale ko. Ye circular fashion mein chalega.

2	3	2	1	5	2	4	5	3	2	3	5	
2°	2°	2'	2'	2°	2'	2°	2°	3°	3°	3'	3'	
		3°	3'	3°	5°	5°	5°	5'	5'	5°	5°	5'
			1°	1°	1°	4°	4°	4°	2	2	2	
x	x	x	x	x	x							



⑤ Enhanced Second Chance Algorithm :

→ Here we consider the modify bit along with the reference bit. With these two bits we can have four possibilities:

(Reference bit, Modify bit)

- | | |
|-----------|--|
| (0 , 0) | → neither recently used nor modified
→ best page to replace |
| (0 , 1) | → not recently used but modified
→ not quite good, as page needs to be written |
| (1 , 0) | → recently used but clean, probably
will be used again soon |
| (1 , 1) | → recently used and modified,
page will be used again & also
need to be written again. |

→ Each page is one of these four classes.

→ The OS goes at most three times searching for the (0,0) class.

① Page with (0,0) ⇒ replace

② Page with (0,1) ⇒ clear modify bit & continue search

③ For pages with reference bit set, the reference bit is cleared.

④ Ek pass kar ke lehiya or koi (0,0) available nahi milta

↳ On 2nd pass a page that was originally (0,1) or (1,0) might have changed to (0,0) ⇒ replace

↳ Agar koi (0,1) ya (1,0) wala hai to modify/clear

↳ By 3rd pass all pages will be at (0,0).

⑥ Additional Bit Algorithm:

→ For each page, OS maintains a group of reference bits 4 or 8 bits. Initially all used bits for each page are set to 0.

→ Every time a page is referenced in a time interval:
↳ Put 1 in place of MSB
↳ then Shift the bits to right after each interval.

→ The page with the lowest reference bits value is the one that is Least Recently Used, thus it will be replaced.

eg:

- 00000000 → The page which is not reference even once
- 11111111 → The page is referenced at least once
- The page x with 11000010⁽¹¹⁹⁹⁾ is more recently used than page y with 01110111 (119)
- ↳ jiski decimal value kam wo LRU.