

HDFS

- ^{distributed} Hadoop File System used to store Big Data.
- HDFS stores metadata on a dedicated server called NameNode / Master Node
- Application data are stored on other servers called DataNodes / Slaves
- All servers are fully connected and communicate with each other using TCP-based protocols.
- DataNodes in HDFS do not use data protection mechanism such as RAID
 - ↳ "Redundant Array of Inexpensive Disk" is a data storage virtualization technology that combines multiple physical disk drive components into one or more logical units for the purpose of data redundancy, performance improvements.

D ARCHITECTURE :

- Files and directories are represented on the NameNode by inodes, which record attributes like permissions, modifications and access times.
- The file content / data is divided into chunks of 128 MB.
- Each block of data is replicated at multiple DataNodes (mostly three nodes) to prevent loss of data

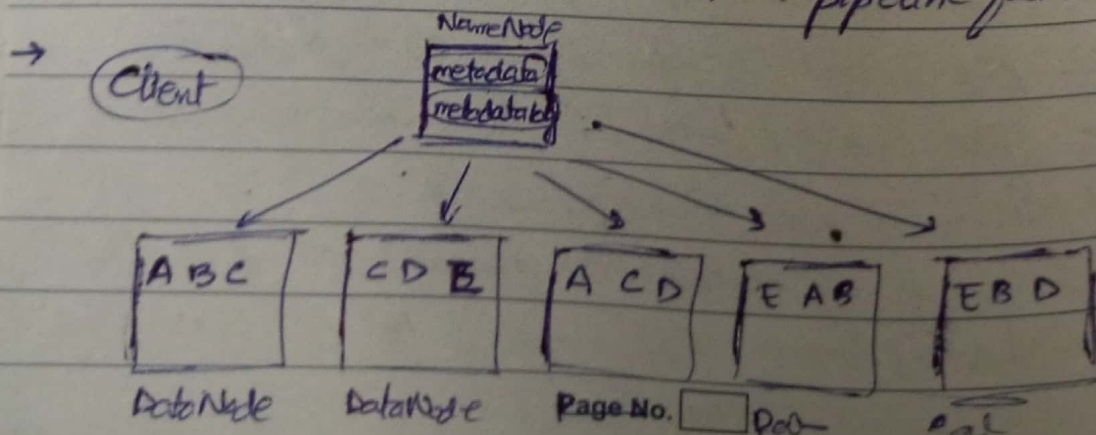
Also stores the modification log of the image called the journal
contains the metadata & list of blocks belonging to each file

Date _____

Day _____

• NameNode :

- Maintains the namespace tree
- Determine the mapping of file blocks to DataNodes (the physical location of data.)
- Stores the metadata
- Collect block reports from DataNodes on block locations.
- Keeps the entire namespace in RAM for fast access
- An HDFS client wanting to read a file it contacts the NameNode for locations of data blocks and then reads the block content from DataNode closest to the client.
- Can have thousands of DataNodes and tens of thousands of HDFS clients / cluster.
- Each DataNode may execute multiple application tasks concurrently.
- When writing the data, the client req the NameNode to nominate a suite of three DN to host the block replicas
- The client then writes to DN in a pipeline fashion.



Heartbeat performs the following tasks: to know total storage capacity, fraction of storage in use and the no. of data transfers currently in progress.

Date _____

Day _____

• DataNodes :

- The DataNodes are responsible for serving read and write requests from the file system's clients.
- The DataNodes also perform block creation, deletion and replication upon instruction from the NN.
- DN periodically send back block reports to NN.
- NN and DN communication is called Heartbeat
- DN send heartbeats to NN to confirm that DN is operating and the block replicas it hosts are available

↳ DN sends Heartbeat after every 3 seconds

↳ Every 10th HB is a block report

↳ NN builds metadata from block reports

↳ If HB not send in 10 mins the NN considers the DN to be out of service

• Block Reports :

- A DataNode identifies block replicas in its possession to the NameNode by sending a block report. A block report contains the block_id, the generation stamp and the length for each block replica the server hosts

- Blockreports provide the NN with an up-to-date view of where block replicas are located on the cluster and NN constructs & maintains latest metadata for blockreports.

Date _____

Day _____

- The 1st BR is sent immediately after the CN registration.
- BR are sent every hour.

• Failure Recovery:

- NN does not directly call DN. It uses replies to heartbeats to send instr. to the DN.

- The instr. include cmds to:

- ① Replicate block to other nodes

- DataNode died

- Copy data to local

- ② Remove local block replicas

- ③ Re-register or to shut down the node

- So when DN died, NN will notice and instruct other DN to replicate data to new DN.

• Image:

- The FsImage is a snapshot of the entire file system metadata, stored as a single file on the disk.
- Contains info about the directories, files, their permissions and block locations.
- The NN loads the image into the memory at startup to reconstruct the file system's metadata.

• Journal :

- A sequential log of all changes made to the file system metadata (eg: file creation, deletion...)
- Tracks update since the last checkpoint.
- If the journal grows too large, it slows down recovery, making checkpoint essential.

• Checkpoint :

- A process where the current metadata (combining the FsImage and edits from the journal) is saved to create a new, consistent image.
- This task is periodically performed by Secondary NameNode.

In order to preserve the checkpoint or journal HDFS can be configured to store the checkpoint and journal in multiple storage directories.

• Startup Process :

When the name node starts :

1. It reads the checkpoint from disk to initialize the namespace image
2. It replays the journal, applying all recorded changes to update the image to its latest state
3. After replaying the journal

- A new checkpoint (update image) is created
- The journal is cleared, starting fresh for new transaction.

• Fault Tolerance :

→ If the checkpoint or journal bcons corrupted or unavailable, the filesystem metadata might be partially/completely lost

→ To prevent this

↳ Multiple storage directories can be configured for check point & journal

↳ Best Practices :

↳ Place directories on different volumes to avoid single-vol failure.

↳ Store one directory on a remote NFS server to protect against node failures.

• CheckpointNode :

→ This node periodically merges the existing checkpoint and journal to create a new checkpoint & clears the journal.

→ How it works :

↳ The CheckpointNode downloads the current checkpoint and journal files from the NameNode

↳ It merges them locally to create an updated checkpoint

↳ The update CP is sent back to NN.

→ Where it Runs :

↳ Usually on a separate host from the NN, ensuring redundancy in case of node failure.

• Backup Node :

→ Like the CN, the BackupNode also creates periodic CP, but it goes further by maintaining an in-memory, up-to-date image of the file system namespace.

→ The image is always synchronized with NN.

→ How it's Better?

↳ The BN doesn't need to download the CP and journal from NN bcz it already has the latest metadata in memory.

↳ This makes CPing faster & more efficient.

→ Act as a read-only NameNode.

↳ Can handle oper. like metadata queries.

↳ It cannot modify the namespace or manage block locations.

→ If NN fails, the BN's image in memory & checkpoint on disk is a record of the latest namespace state.

Date _____

Day _____

→ HDFS snapshots are created to safeguard the file system during software upgrades, reducing the risk of data corruption caused by bugs or human error. Snapshots persistently save the current state of both data & metadata, allowing administrators to rollback to the pre-upgraded state if issue arises, ensuring data integrity & recovery.

• FILE I/O OPERATIONS :

HDFS follows a "single writer, multiple readers" model.

• File writing & Lease :

- A client writing to a file is granted a lease, ensuring exclusive write access.
- The lease is periodically renewed via a heartbeat sent to NN.
- Once the file is closed, the lease is revoked, and no further modifications can be made (except appending).

• Lease Limits :

↳ Soft Limits : Ensures exclusive access for a set duration; after this, another client may preempt the lease if the writer fails to renew it.

Means another client can preempt the lease.

also

soft

hard Day

↳ Hard Limits : After one hour without renewal, HDFS assumes that writer has quit, closes the file, and recovers (take back) the lease.

↳ The new client can only request for lease after soft limit expiry & after hard limit expiry only then he can write.

• Concurrent Reads :

The writer's lease does not block other clients from reading the file, allowing many readers to access it simultaneously.

1. File Write Process :

1. Block Allocation :

- The NN allocates a new block with a unique ID and selects a list of DN to host its replicas.
- The DN form a pipeline to minimize network distance.

2. Data Transmission :

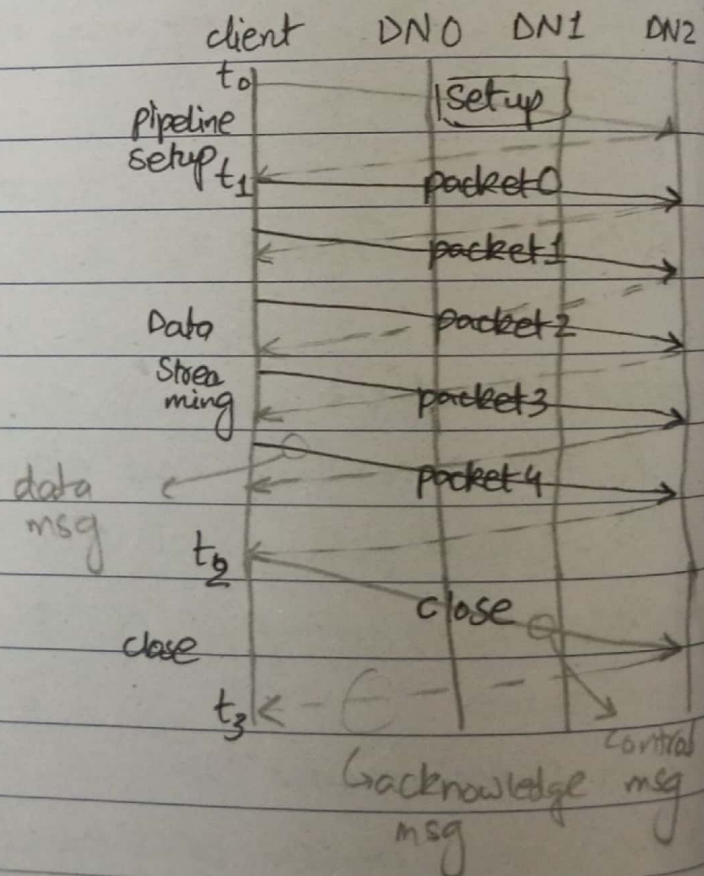
- The client buffers data locally.
- Once the buffer (typically 64 KB) is full, data is split into packets and pushed into the pipeline.
- Bytes are pushed into the pipeline as a sequence of packets.

Page No.

Date _____

Day _____

- After a packet buffer is filled, the data are pushed to the pipeline
- The next packet is pushed, after receiving the acknow. from prev one.
- 3. Visibility of data:
 - Data is not guaranteed to be visible to new reader until the file is closed.
 - The flush operation ensures visibility by waiting for acknowledgments from all DN in the pipeline.
- 4. Pipeline Stages:



2. File Read Process:

1. A client req the block location from the NN.
2. The client reads data directly from one of the DN hosting the block.

Page No.

3. Checksums

↳ HDFS verifies block integrity using checksums stored on the DataNode

↳ If corruption is detected, the client notifies the NN & fetches an alternate replica.

3. Data Integrity With Checksums:

1. Checksum generation:

- When writing data, the client computes checksums for each block & sends them to DN along with the data.
- DN stores these checksums in a separate metadata file.

2. Checksum verification:

- During reads the client computes the checksum for the data it receives and compares it with stored one
- If there is a mismatch, it detects corruption

3. Corruption Handling:

- The client inform NN about the corrupt replica
- The client then goes to differ replica.

Date _____

Day _____

▷ REPLICATION MANAGEMENT :

1. Block Placement :

- **Cluster Topology :** Nodes are grouped into racks connected via switches, with faster communication within racks than b/w them.
- **Rack Awareness :**
 - HDFS estimates network bandwidth based on distance b/w nodes
 - A rack identification script helps the NN determine a DN's rack
- **Default Placement Policy :**
 - First replica on the writer's node
 - 2nd & 3rd replicas on two nodes in a diff rack
 - No DN hosts more than one replica of a block
 - No rack has more than two replicas ~~unless~~ provided there are sufficient racks on the cluster
- **Pipeline and Proximity :**
 - Data is written in a pipeline order
 - Blocks are read based on the proximity to the client.

2. Replication Management :

- Under-replication :
 - The NN adds under-sep blocks to a priority queue
 - Blocks with a single replica are the highest priority for replication
- Over-replication :
 - Extra replicas are removed, preferring to :
 - Keep replicas on diff racks
 - Remove from DN with least available space
- Objective : Maintain balanced storage while ensuring data availability.

3. Balancer : (An Application Program)

- Purpose : Balances disk space utilization across DN without reducing replicas or rack diversity.
- Operation :
 - Moves replicas from highly utilized nodes to underutilized nodes
 - Ensures no impact on data availability during balancing.

Date _____

Day _____

4. Block Scanner :

- Purpose: Periodically runs on each DN and block data integrity using checksums.
- Operation:
 - Adjust read bandwidth to complete scans in a set time
 - Marks corrupt blocks and informs the NN.
 - Corrupt replicas are flagged but not deleted immediately.

5. Decommissioning :

- Process :
 1. Cluster administrator marks a DN for decommissioning using an exclude list.
 2. The Data Node
 - Stops receiving new replicas
 - Continues serving read request
 3. The NN replicates its blocks to other nodes
 4. Once all blocks are replicated, the DN enters decommissioned state and can be removed safely.

Date

"PRACTICE AT YAHOO"

Day

▷ Durability of Data :

◦ Uncorrelated Node Failure :

Replication of data three times is a robust guard against loss of data due to uncorrelated node failures.

◦ Correlated Node Failures :

The failure of a rack or a core switch, HDFS can tolerate losing a rack switch (each block has a replica on some other rack because no more than two replicas are stored on same rack.)

◦ Loss of electrical Power to Cluster :

A large cluster will lose a handful of blocks during a power-on restart.

▷ Automated Failover :

↳ Plan :

↳ Zookeeper, Yahoo's distributed consensus technology to build an automated failover solution.

• Scalability of the NameNode.

↳ Solution :

↳ Our near-term sol to scalability is to allow multiple namespaces (and NameNodes) to share the physical storage within a cluster.

Page No.

NOTE

Date _____

Day _____

↳ Drawbacks :

↳ The main drawback of multiple indep name space is the cost of managing them.

▷ ADVANTAGES OF HDFS

① Load Balancing: As data is distributed, load on datanodes is also distributed. This allows improved performance.

② High Availability: By default HDFS replicates 3 copies of each block. This ensures fault tolerance and high availability.

③ Data Localization: Replicated copies of data facilitates localized access of data, which reduces network latency.

• But HDFS entails additional cost of data storage.