

Date _____

Day _____

"CHAPTER # 05"

• Error Checking:

→ If compiler doesn't support the openmp then we should 1st check.

```
#ifdef _OPENMP  
# include <omp.h>  
#endif
```

1

```
int main() {
```

```
#ifdef _OPENMP
```

```
    int my_rank = omp_get_thread_num();
```

```
    int thread_count = omp_get_num_threads();
```

```
#else
```

```
    int my_rank = 0;
```

```
    int thread_count = 1;
```

```
#endif
```

1

3

Date _____

Day WED

▷ Reduction :

```
global_result = 0; // shared variable  
#pragma omp parallel num_threads(thread_cnt)  
{  
    double my_result = 0.0; // private variable  
    my_result += Local_trapezoid(double a, double b, int n);  
}
```

#pragma omp critical \rightarrow multiple threads access same
resource to once condition.

→ Now using reduction

```
global_result = 0;  
#pragma omp parallel num_threads(thread_cnt) \  
reduction (+: global_result)  
global_result += Local_trapezoid();
```

→ Use of subtraction operator (Problematic!)

result = 0;

```
for (i=1; i<=4; i++)
```

result -= i;

↳ result stores -10. →

Date _____

Day _____

But if this loop was performed by two threads
then iterations would have been divided.

$$T_1 \rightarrow \text{result} = -3 \quad (-1-2)$$

$$T_2 \rightarrow \text{result} = -7 \quad (-3-4)$$

And then subtracting the results of these two
would give 4 ($-3 - (-7) = 4$)

which is not correct it should have been -10
($-3 + (-7) = -10$) .

→ Use of float or double as reduction variable
is also problematic.

As in floats $(a+b)+c$ may not be exactly
equal to $a+(b+c)$.

▷ Parallel for :

→ OpenMP can only parallelize the for loops whose
no. of iterations are known.

```
for ( ; ; )      or    for (i=0; i<n; i++) {  
}                                                            if (—)  
                                                              break;  
                                                              }  
                                                              }
```

These cannot be parallelized.

→ While & Do-while can also be not parallelized.

Date _____

Day _____

⇒ Data Dependencies :

$$\text{fib}[0] = \text{fib}[1] = 1 ;$$

```
#pragma omp parallel for num_threads(threadt)
```

```
for (i=2; i<n; i++)
```

$$\text{fib}[i] = \text{fib}[i-1] + \text{fib}[i-2] ;$$

→ Here the result of $\text{fib}[2]$ is dependent on $\text{fib}[1]$ & $\text{fib}[0]$ therefore there is dependency among loops.

→ OpenMP doesn't check for dependencies among iteration in a loop

→ A data dependent loop gives unpredictable results when aligned by OpenMP.

→ These type of data dependency also called loop-carried dependence.

↳ result of one iteration depends on result of another iteration.

for (i=0; i<n; i++) { → yaha par general dependence hai.

$x[i] = a + i * h ;$ loop carried nhii. loop carried main previous iteration par

$y[i] = \exp(x[i]);$ dependence hoti hai.

} like $n[i] = n[i-1] + i^2 * h ;$

Date _____

► Estimating π :

$$\pi = 4 [1 - \frac{1}{3} + \frac{1}{5} - \frac{1}{7} + \dots] = 4 \sum_{k=0}^{\infty} \frac{(-1)^k}{2k+1}$$

double factor = 1

double sum = 0

for ($k = 0$; $k < n$; $k++$) {

sum += factor / (2 * k + 1);

} factor = -factor; → This factor value changed in
 one iteration will be used
 by next iteration.

pi-approx = 4.0 * sum;

So loop carried dependency

& this should be a
private variable or else
if it is shared then it would
cause race condition.

► Corrected Version

double sum = 0.0;

pragma omp parallel for reduction(+:sum) \\\
private(factor)for ($k = 0$; $k < n$; $k++$) { if ($k \% 2 == 0$) → isse ye previous iterations

factor = 1.0; par dependent whi hoga

 like 'k' ki current value
 par dependent shega.

else

factor = -1;

sum += factor / (2 * k + 1);

}

Page No.

pi-approx = 4.0 * sum,

▷ Bubble Sort - Loop Carried Dependencies

```
for (listLength = n; listLength >= 2; listLength--) {
```

```
    for (i = 0; i < listLength - 1; i++) {
```

```
        if (a[i] > a[i + 1]) {
```

```
            swap(a[i], a[i + 1]);
```

```
}
```

```
}
```

```
}
```

Outer Loop Dependence:

- ① Each pass of outer loop depends on completion of previous passes.

Example:

If list starts with [3, 4, 2, 1], after pass-1 list becomes [3, 1, 2, 4]. The 2nd pass should operate on [3, 1, 2, 4].

Now if 1st & 2nd pass exec simultaneously, the 2nd pass could operate on [3, 4, 2, 1].

Inner Loop Dependence:

- ② Each inner loop iteration compares & possibly swaps elements, impacting next iteration.

Date _____

Example:

In a list [5, 1, 2], starting with $i=0$, we compare $a[0] > a[1]$ & swaps ($a[0], a[1]$) i.e. 3 8 1. [1, 3, 2]. The next comparison when $i=1$ $a[1] > a[2]$ & swaps (3 & 2).

[1, 2, 3].

Now if $i=0$ & $i=1$ iteration runs || by then $i=1$ could incorrectly compare 1 & 2 before $i=0$ completes it's swap.

Odd-Even Sort Parallel Program:

```
# pragma omp parallel num_threads(thread_count) \n\n    default(private) shared(a,n) private(i,phase)
```

```
for (phase=0; phase<n; phase++) {\n    #pragma omp parallel for           // in even phase compare\n    for (i=1; i<n; i+=2) {            // the odd subsubts\n        if (a[i-1] > a[i])           // with left elements\n            swap(a[i-1], a[i]);\n    }\n}
```

```
else                                // in odd phase compare odd sub\n    #pragma omp parallel for          // to right element\n    for (i=1; i<n-1; i+=2) {\n        if (a[i] > a[i+1])\n            swap(a[i], a[i+1]);\n    }
```

2 4 3

► Schedule Clause :

Schedule (type, chunkSize)

→ TYPES OF SCHEDULE CLAUSE :

① STATIC SCHEDULE TYPE :

The system assigns chunks of chunk-size iterations to each thread in a round-robin fashion.

If we have 12 iteration then each thread will get

$T_0 : 0, 3, 6, 9$

$T_1 : 1, 4, 7, 10$

$T_2 : 2, 5, 8, 11$

If schedule (static, 2)

$T_0 : 0, 1, 6, 7$

$T_1 : 2, 3, 8, 9$

$T_2 : 4, 5, 10, 11$

If chunk-size is omitted, then it is approximately equal total iterations / thread count.

② DYNAMIC AND GUIDED TYPE :

In a dynamic scheduling the iterations are broken up into chunks of chunk-size. When one thread finishes a chunk, it request for another.

Date

Benefits:

- ↳ Dynamic scheduling is effective for loops with variable execution times for each iteration.
- ↳ Faster threads don't need to wait for slower threads, as they can pick up new chunks when available.

In Guided Scheduling in OpenMP is a type of dynamic scheduling where the size of the chunks decreases as more iterations are assigned to the threads.

As chunks are completed, the size of subsequent chunks decs, based on

$$\text{chunk_size} = \frac{\text{remaining iterations}}{\text{no. of threads}}$$

Benefits:

- ↳ Larger chunks at beginning help handle initial heavy workloads
- ↳ Useful when workload / iteration reduces gradually throughout the loop.

`schedule(guided, 2)` → tells OpenMP to start with larger chunks, progressively dec. the chunk size until reaches 2.

Date _____ 1 Day _____

③ RUNTIME TYPE :

When we write `schedule(runtime)`, the system uses the environment variable `OMP_SCHEDULE` to determine at runtime how to schedule loops.

We can set the value of env variable by running following

```
$ export OMP_SCHEDULE = "static,1"
```

Now when we run our program the system takes will take the value from `OMP_SCHEDULE` and our clause will be modified to

```
schedule(static,1)
```

→ PRODUCER AND CONSUMER

→ CACHE, CACHE COHERENCE AND FALSE SHARING :

→ The design of cache takes into consideration the principles of temporal & spatial locality.

"if a processor access a main memory location x at a time 't', then it is likely that at times close to 't', it will access main memory locations close to 'x'."

→ Therefore instead of transferring only ' x ' into cache a block of memory is transferred called cache line or cache block.

Page No.

Thread Safe :

- A function is not thread safe if it give different result when executed by multiple threads.

→ `#pragma omp critical
y = f(x);`

`double f(double x) {`

`#pragma omp critical
z = g(x)`

}

This will create deadlock

If Thread 1 started executing the 1st critical section then

no thread can enter in any critical section so, we call the function

Again & start to execute it has a nested critical section & it cannot enter into it.

It will wait infinitely.

→ We can differentiate the CS by giving names to them

`#pragma omp critical (one)
y = f(x)`

`double f(double x) {`

`#pragma omp critical (two)`

`z = g(x);`

"PRODUCE CONSUMER"

► Message Passing :

- Threads communicate each other via messages
- Threads can access same memory space, which allows them to share info.
- Each thread has its private queue
- If one thread want to send msg to another it places msg in other's queue (destination queue).
- Receiving thread checks its queue.
- If there is a msg, the thread dequeues it & process it.

```
for (send-msg=0 ; send-msg < max ; send-msg++) {
```

```
    send-msg(); // Each thread sends a max
```

```
    Tryreceive-msg(); // amount of msg to other  
                      // thread & checks if it  
                      // delivers any msg
```

```
}
```

```
while (!Done()) { // If all msgs are send but
```

```
    Try-receive(); // still receiving the msg  
                  // from other threads.
```

Date

▷ Sending Message

- If two threads tries to simultaneously enqueue a msg in a destination thread's queue - race condition may occur.

Send-msg() {

msg = random();

dest = random % thread-count;

#pragma omp critical

Enqueue(queue, dest, my-sample, msg);

▷ Receiving Message

- Only the thread to which the queue belongs can dequeue the msg.

→ We uses enqueue variable & dequeue variable to keep the track of no. of msgs enqueued & dequeued.

→ This is done so that we can get queue-size.

$$\text{queue_size} = \text{enqueue} - \text{dequeue}$$

- No conflicts occur when we simultaneously enqueue or dequeue msg.

Date _____

Day _____

▷ Termination Detection :

- If owner thread execute
Done () {

- 1 queue-size = enqueue - dequeue
- 2 if (size == 0) return TRUE
- 3 else return FALSE.

3

If a thread 'U' executes line 1 and at same time
thread 'V' sends a msg then thread U can't receive
it as it will return ~~size~~ size as queue-size == 0.
To avoid this we do ^{TRUE}

```
if (size == 0 && done_sending == thread_count)
    return TRUE;
else
    return FALSE;
```

▷ Startups:

- Master thread get command line arguments &
allocate array of msg queue to each thread.
- If a thread receives a msg before completing
its memory allocation for queue it would
crash the program.
- We should create an explicit barrier so that
no thread could send msgs before complete allocation
of memory.

Date _____

Day _____

- Atomic directive can only protect critical section that consist of a single C assignment statement

$x++$

$--x$

$x = (y++)$; ^{→ this remains unprotected.}

④ Locks:

`omp_lock_t lock;`

`omp_init_lock(&lock);`

`omp_set_lock(&lock);`

`omp_unset_lock(&lock);`

`omp_destroy_lock(&lock);`

"OpenMP Parallel Programming"

• Parallel Region :

```
#include <stdio.h>
```

```
#include <omp.h>
```

```
void main() {
```

```
    omp_set_num_threads(4); // This will create 4
```

```
    int id;
```

threads. If we don't
write this default=2.

```
#pragma omp parallel alike andas seasa code 4  
bars chalega // by -
```

```
{
```

```
    id = omp_get_thread_num();
```

```
    printf("Thread %d : HelloWorld\n", id);
```

```
}
```

```
return 0;
```

```
}
```

```
gcc -fopenmp hellocool.c  
• a.out
```

Date _____

Day _____

▷ Parallel For Loop:

#pragma omp parallel for
for(—————) {

⇒ This loop will execute in ||.

}

or

#pragma omp parallel
{

#pragma omp for
for(—————) {

}

}

#pragma omp parallel
for(—————) {

{

} → Is toha kaise to loop ki iteration a
nhi hongi bike ka thread mein saare
iterations chalengi.

#include <stdio.h>

#include <omp.h>

void main() {

 int i, sum = 0;

Infinix NOTE

Date _____

Day _____

```
int thread_sum[4];  
omp_set_num_threads(4);
```

```
#pragma omp parallel  
{
```

```
    int ID = omp_get_thread_num();  
    thread_sum[ID] = 0;
```

ye loop ki iteration khel
in 25, 25 mein divide ho
jaglegi.

```
#pragma omp for  
for(int i=1 ; i<=100 ; i++)  
    thread_sum[ID] += i;
```

7

```
for(i=0 ; i<4 ; i++)
```

```
    sum += thread_sum[i];
```

8

▷ Shared & Private Variables.

All variables declared outside || region are shared among all.

Variables inside || region are private to each thread. Each thread has a local copy for it.

Iteration variable are by default private.
Shared: sum, thread_sum[], private: i, ID

value of private variable is unspecified at beginning of
ll block & also after completion of ll block

Date _____

Day _____

Explicit rule:

#pragma omp parallel for shared(n, a)
for (int i=0; i<n; i++)

int b = a + i;

*

int a, b, c, n

#pragma omp parallel for default(shared) private(a, b)
for () {

! ll a & b are private

! ll c and n are shared

int n = 10;

std::vector<int> vector(n);

int a = 10;

ye ikhne ki wajah se humein
shared waale explicitly btaane
pasenge.

#pragma omp parallel for default(none) shared(n, vector, a)
for ()

{ }

④ Critical Sections:

#include <omp.h>

int count = 0;

void myFunc() {

#pragma omp critical

Day D Barrier :

It is the point in the execution of a program where threads wait for each other.

No thread is allowed to continue until all threads reach barrier.

Used to achieve sync.

As soon as one thread reaches barrier then all threads in the team must reach the barrier

```
#pragma omp barrier
```

OpenMp function

① void omp_set_num_threads(4)

int omp_get_num_threads() → set no. of threads in current team

int omp_get_max_threads() → Ret max no. of threads that could be used to form a new team using a `ll` construct without a `num_threads` clause.

int omp_get_thread_num() → Ret ID of current thread

int omp_get_num_procs() → set no. of processors available to the program

Date _____

`int cmp_in_parallel()` → Ret true if the call to routine
is enclosed by an active `||`
region

▷ Parallel For Loops : Scheduling

→ If working of each thread increase after each iteration then the last thread will have lot of work to do.

→ Therefore we need to distribute / load balance the work.

```
#pragma omp parallel for schedule(static, 1)
for(int i=0; i<16; i++)
    c(i);
```

3

| | | | | |
|------------------|------|------|-------|-------|
| T ₀ : | w(0) | w(4) | w(8) | w(12) |
| T ₁ : | w(1) | w(5) | w(9) | w(13) |
| T ₂ : | w(2) | w(6) | w(10) | w(14) |
| T ₃ : | w(3) | w(7) | w(11) | w(15) |

"1" indicates that we assign 1 iteration to each thread before switching to next thread.

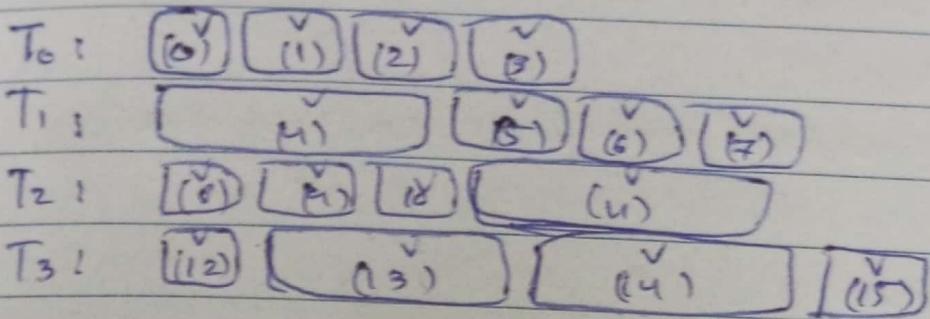
Date _____

Day _____

• Dynamic Scheduling:

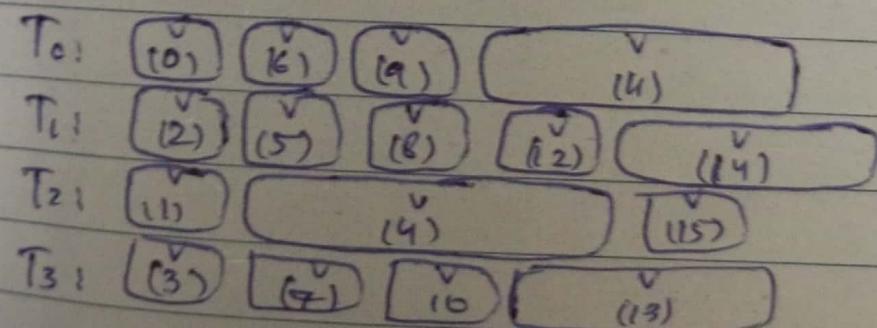
→ Kisi Kabar kuch threads ~~toh~~ boht jaldi execution khatam kareti hai or kuch ~~toh~~ boht der se finish hoti hai.

⇒



→ Iteration k chunk size hai.

```
#pragma omp parallel for schedule(dynamic, 1)
for(int i=0; i<16; i++)
    v(i)
```



Dynamic Scheduling assigns iterations of loop to threads at runtime, distributing a specified chunk of iterations to each thread. After a thread finishes its chunk, it gets a new one allowing for better load balancing, especially when iterations have varying execution time.

Date _____

④ Reduction:

- ↳ Reduction is used to combine variables from diff threads into a single result after execution.
- ↳ Specifying operator, defines how individual thread results will be combined (e.g., + for sum, * for product).
- ↳ This ensures that each thread work on a private copy of the variable, avoiding direct access to shared mem during computation.

```
int sum = 0;
```

```
#pragma omp parallel for reduction(+:sum)
```

```
for (int i = 1; i <= 10; i++)
```

```
    sum += i;
```

```
printf("Sum = %d", sum);
```

```
return 0;
```

Precche example mein jo humne thread_sum[] ke array bnaya phir har thread ke liye use use krrao or last mein combine krrao ye soora kaam reduction ki deta hai.

Date _____

Day _____

► Sections:

- ↳ Section is used to divide a program into independent code blocks, each of which can be executed by different threads concurrently.
- ↳ It allows // execution of non-loop tasks.

```
#pragma omp parallel sections num_threads(3)
```

{

```
#pragma omp section
```

{

```
    printf("Task by threads %d", omp_get_thread_num());
```

}

```
#pragma omp section
```

{

```
    printf("Hello world");
```

}

```
#pragma omp section
```

{

```
    printf("My World");
```

}

G