

CHAPTER # 07

- Transport Layer -

- TCP provides logical communication b/w application processes running on different hosts.
- TCP is implemented at "end-systems", not in network routers.
 - ↳ TCP on sender side → breaks application messages into segments, passes to network layer
 - ↳ TCP on receiver side → reassembles segments into msgs, passes to application layer.

- ▷ Relationship Between Transport and Network layer
- Transport layer is communication b/w processes.
- Network layer is communication b/w hosts
- Within an end system, a transport layer protocol moves msgs from app processes to network edge (network layer) & vice versa.
- Agar koi eski service like bandwidth or delay guarantee jo network layers provide nahi kr sakta to TCP transport layer protocol bhi nahi kr sakta.
- Lekin kuch services like reliable data transfer, encryption jo TCP provide karta hai tekn NL nahi karta.

Date

Day

→ Transport layer actions

▷ Sender :

- ① TL is passed an application layer message
- ② determines segment header values
- ③ create segment
- ④ passes segment to NL

▷ Receiver :

- ① Receiver segment from NL
- ② checks header values
- ③ Extract application layer msgs
- ④ demultiplexes message up to application via socket

• TCP (short intro)

↳ reliable, in-order delivery

↳ congestion control

↳ flow control

↳ connection setup

• UDP

↳ unreliable, unordered delivery

↳ "best-effort" [it makes best effort to deliver segments but does not guarantee delivery services]

Date _____

This occurs on the receiver's ability to receive the sender's packet.
due to buffer overflow (occurs on intermediate nodes) → occurs on receiver Day

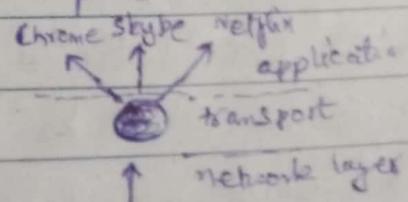
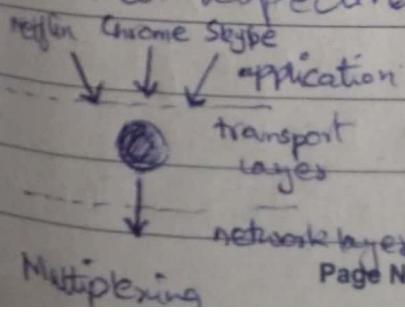
Congestion Control	Flow Control
① Prevents network-wide congestion (overloaded routers)	① Prevents sender from overwhelming the receiver
② Focused on entire network traffic	② Focused on individual connection b/w sender & receiver
③ Applied when network is overloaded (packets get dropped)	③ Applied when receiver can't handle incoming data fast enough
④ Reduces sender's rate if congestion is detected	④ Receiver advertises a window size (how much it can handle) - Sender adjust its sending rate accordingly.

o Multiplexing & Demultiplexing :

→ Multiplexing - gather / collect data from multiple sockets (processes communicate with TLP using sockets)

, add transport header (like src, destination port nos, ip addresses etc) & send to NL

→ Demultiplexing - receives segment from NL , uses the header info & deliver the segment to their respective sockets (process).



→ IP addr of src & dest is in IP datagram's header.
Date _____ Day _____

→ Each socket is associated with a PORT number ranging from 1024 - 65535. (16-bit numbers)

→ The header of a segment contains 32-bit src & dest PORT numbers.

▷ Connection less Multi/Demultiplexing

▷ How Demultiplexing work? (UDP)

① Host receives IP datagram

i- datagram has src & dest IP

ii- each datagram carries one TLP segment

segment

iii- each segment has src & dest PORT

② host uses IP addr & PORT number to direct segment to appropriate socket

→ UDP socket is identified by 2-tuple (destination IP, dest PORT)

→ A datagram with same dest. PORT #, but different src. IP address and/or source PORT will be directed to "same socket" at receiving host

▷ Connection Oriented Multi/Demultiplexing (TCP)

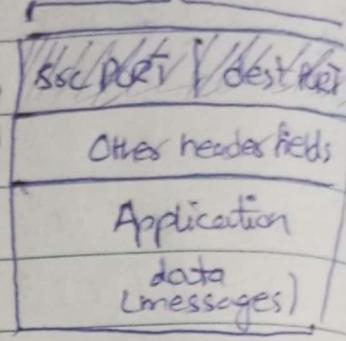
→ TCP socket is identified by a 4-tuple (src IP, src PORT, dest IP, dest PORT)

→ The host (destination node) uses this tuple for demultiplexing

→ Kisi ka socket ka tuple unique hogा होमेहा.

→ Server may support many simultaneous TCP sockets

→ Each socket associated with Page No. a different connecting client.



clientSocket = socket (AF_INET, SOCK_STREAM)

clientSocket.connect (serverIP, (2000)) → dest PORT#

→ TCP client creates a socket & sends a connection establishment req. → client sends a TCP segment with dest PORT# & a special connection bit set to 1 in TCP header. Segm also include src PORT#

At Server:

connectionSocket, addr = serverSocket.accept()

→ The server receives request & saves the 4-tuple value.

▷ Web Servers:

→ Web server main task process to lie along socket which create hole (or it will seriously effect performance)

→ Web servers often use only one socket process, and create a new thread with a new connection socket for each new client connection

→ If client & server are using persistent HTTP, then throughout the duration of persistent connection the client & server exchange HTTP msgs via the same server socket.

→ For non-persistent HTTP for each request/response a new socket is created (a new TCP connection is created).

→ happens at all layers

"Multiplexing is based on segment & datagram header field values"

► [USER] DATAGRAM Protocol :

- It is very simple transport protocol.
- Apart of multi/demultiplexing & some little error checking it do nothing
- Segments may be "lost" → "delivered out of order"
- DNS uses UDP.
 - ↳ Constructs a DNS query & passes to UDP
 - ↳ Without any handshaking, UDP adds header field to message & passes to network layer
 - ↳ NL encapsulates UDP segment into datagram & sends to name server
 - ↳ The host DNS waits for reply

- Finer application control over what data is sent, and when :

↳ UDP sends data immediately without waiting for congestion control or acknowledgments

↳ Realtime application need low latency & can tolerate some packet loss

- No connection establishment :

↳ UDP does not require a 3-way handshake, unlike TCP

↳ Faster transmission i.e why DNS uses UDP to send domain query names quickly -

- No connection state : (NO RTT req)
 - ↳ TCP maintains buffers & control parameters,
UDP doesn't
- Lower Overhead :
 - ↳ TCP header = 20 bytes
 - ↳ UDP header = 8 bytes

→ HTTP now also runs on UDP, providing their own errors & congestion control at application layer

- UDP segment Structure →

src PORT	dest PORT
length header+data	checksum
Application data	
(message)	

- Check Sum :

1110011001100110
 1101010101010101
1011101110111011
 ↴
 sum
 → 1011101110111000 → Sum
 ↴ complement
 ↴ checksum.

→ but there is a flaw in this as.

$$(\text{sender side}) \quad 4 + 6 = 10$$

$$(\text{receiving end}) \quad 7 + 3 = 10 \rightarrow \text{but no error}$$

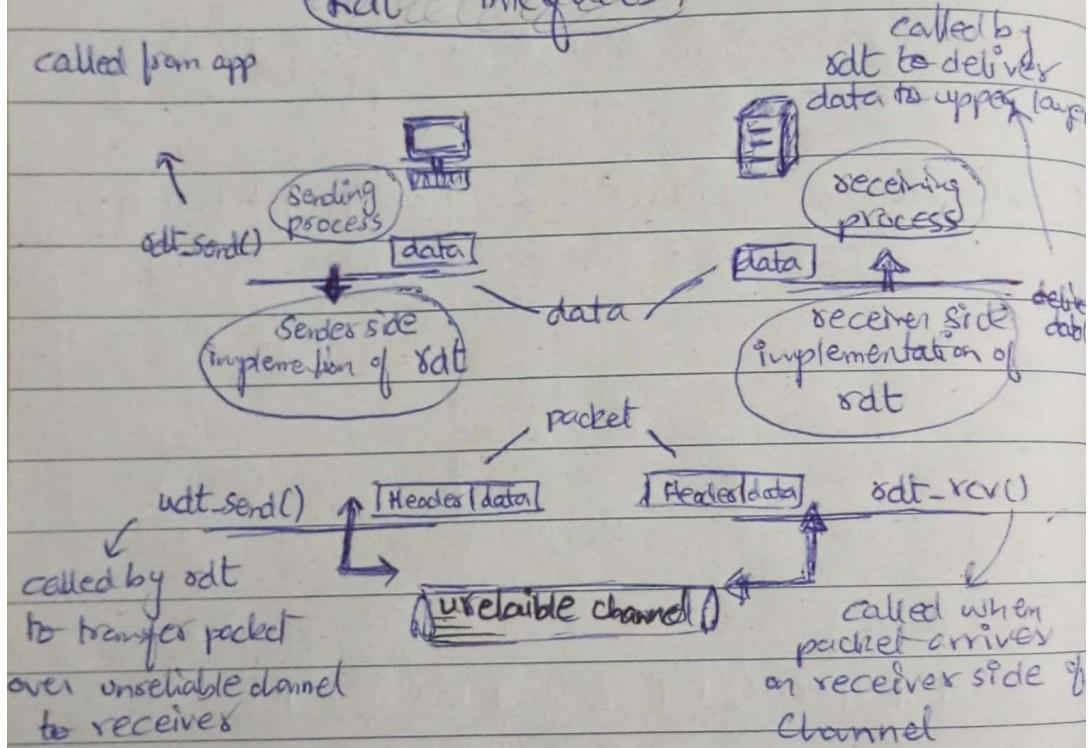
detected as
sum is same

changed the
data but sum is
same.

► PRINCIPLES OF RELIABLE DATA TRANSFER:

Rdt Interfaces

called from app



① rdt 1.0 : Reliable Data Transfer Over a Reliable Channel

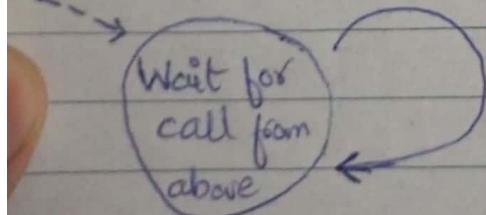
the event causing the transition

rdt_send(data)

packet = make_pkt(data)

udt_send(packet)

↳ action taken when transition occurs



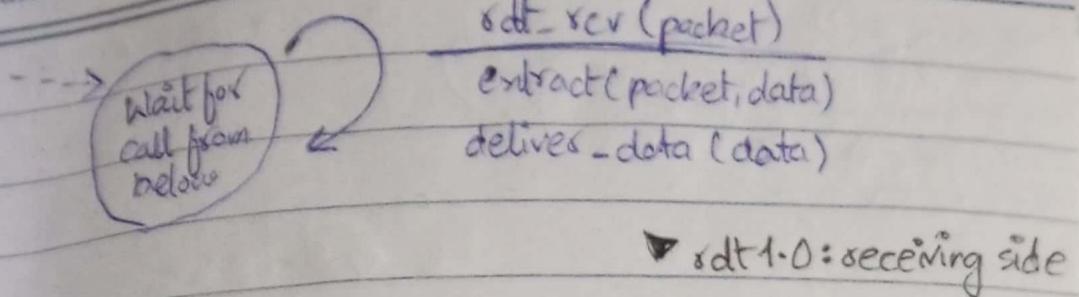
↑ → this symbol is used to represent no action / no event

Page No.

► rdt 1.0 : sending side

Day

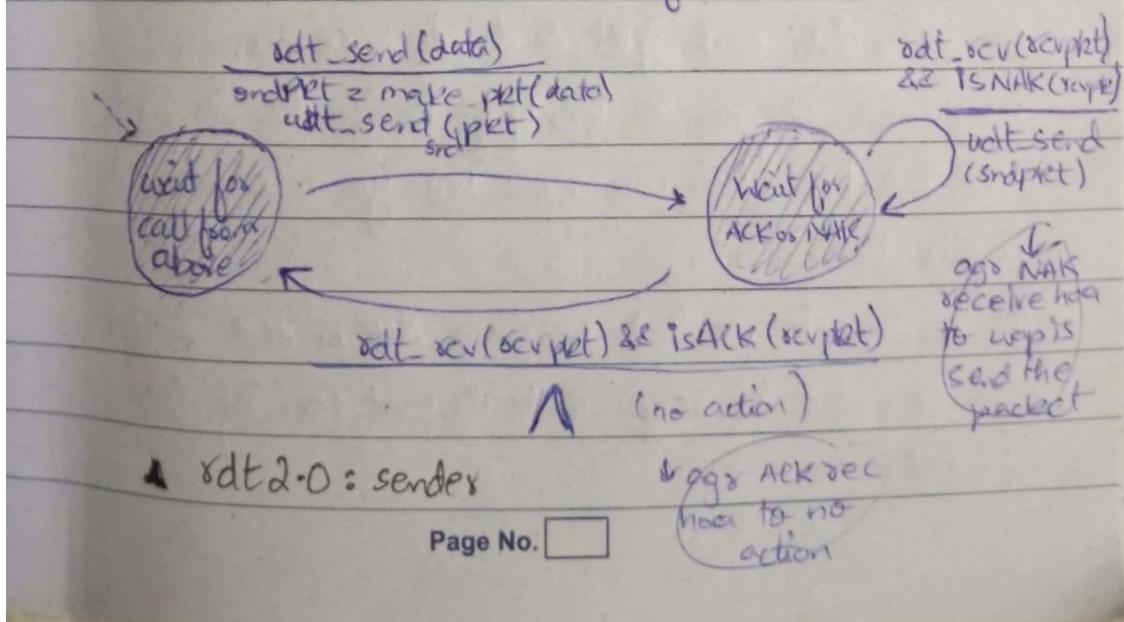
Date



► rdt 1.0 : receiving side

- rdt 2.0 : Channel with bit errors
 - ↳ underlying channel may flip bits in packet
 - ↳ checksum to detect bit errors.
- transmission based on ACK & NAK is called ARQ protocols (Automatic Repeat reQuest)
- acknowledgments (ACKs) : receiver explicitly tells sender that packet received is OK
- neg acknowledgment (NAKs) : receiver explicitly tells sender that packet had errors.
- sender retransmits pkt on receipt of NAK

"Stop & wait" \Rightarrow sender sends one pkt, then waits for receiver's response



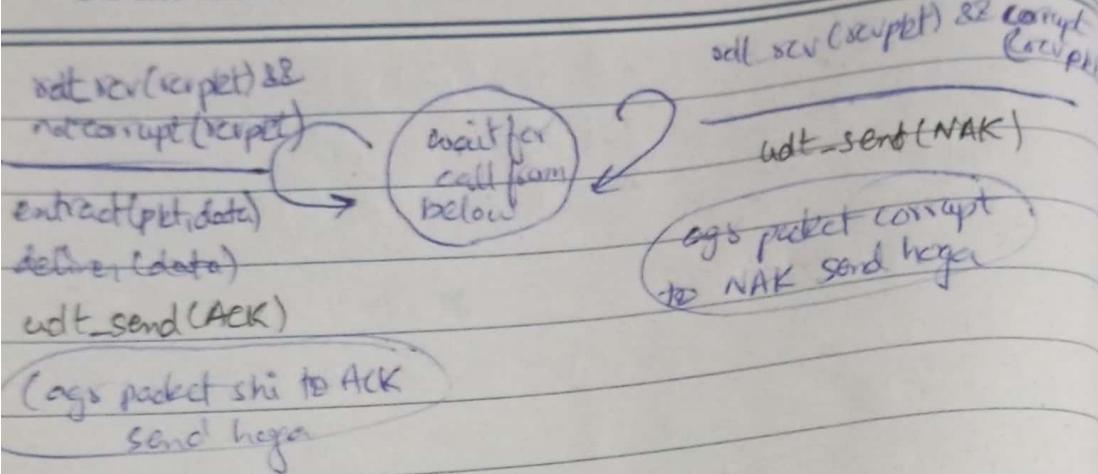
► rdt 2.0 : sender

Page No.

→ go to ACK rec
block to no
action

Date

Day



→ rdt 2.0 has a fatal flaw:

↳ ACK or NAK can also be corrupted

↳ How can then sender know that receiver received the response or not?

↳ we can retransmit data if ACK or NAK is corrupted

↳ Agar humein sec ki taraf se ACK receive ho jata tha or ac corrupt hogya to sender wapis send krdEGA lekin phir sec ke poori duplication ho jayegi data ka.

this is a 1 bit no. (0081)

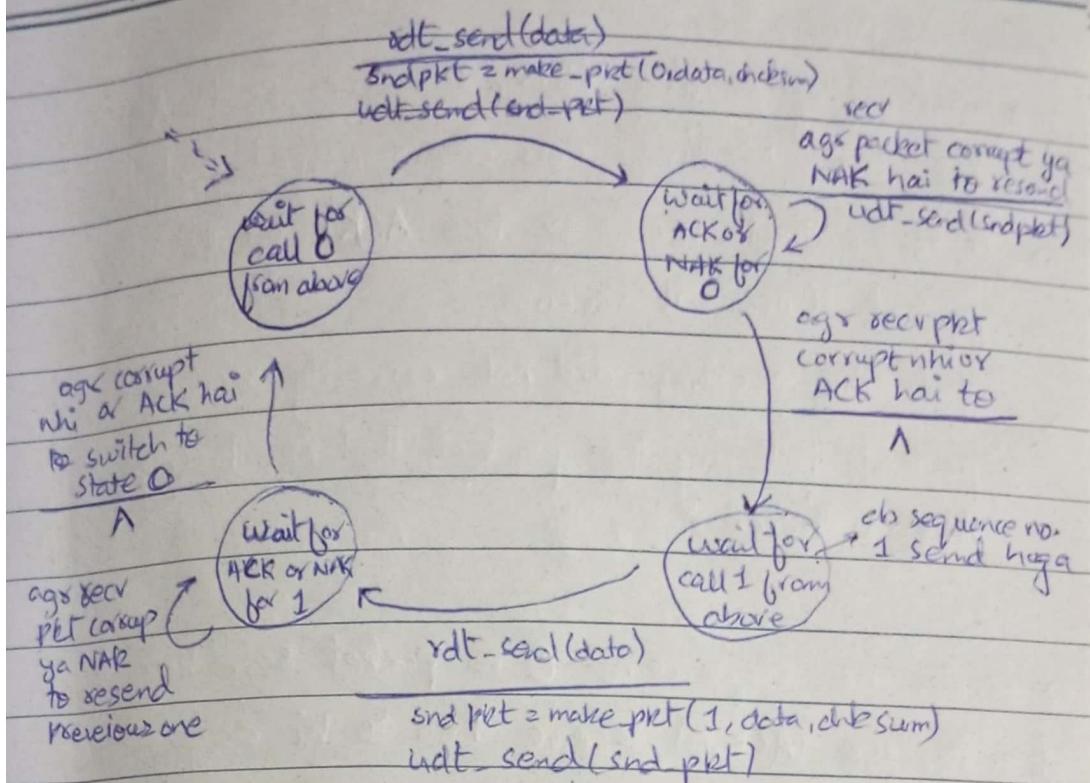
To deal with this problem we add sequence number with each pkt send to receiver.

↳ receiver seq.no se identify krega ko jo pkt receive hoa hou wo new hou ya retransmitted wala hou.

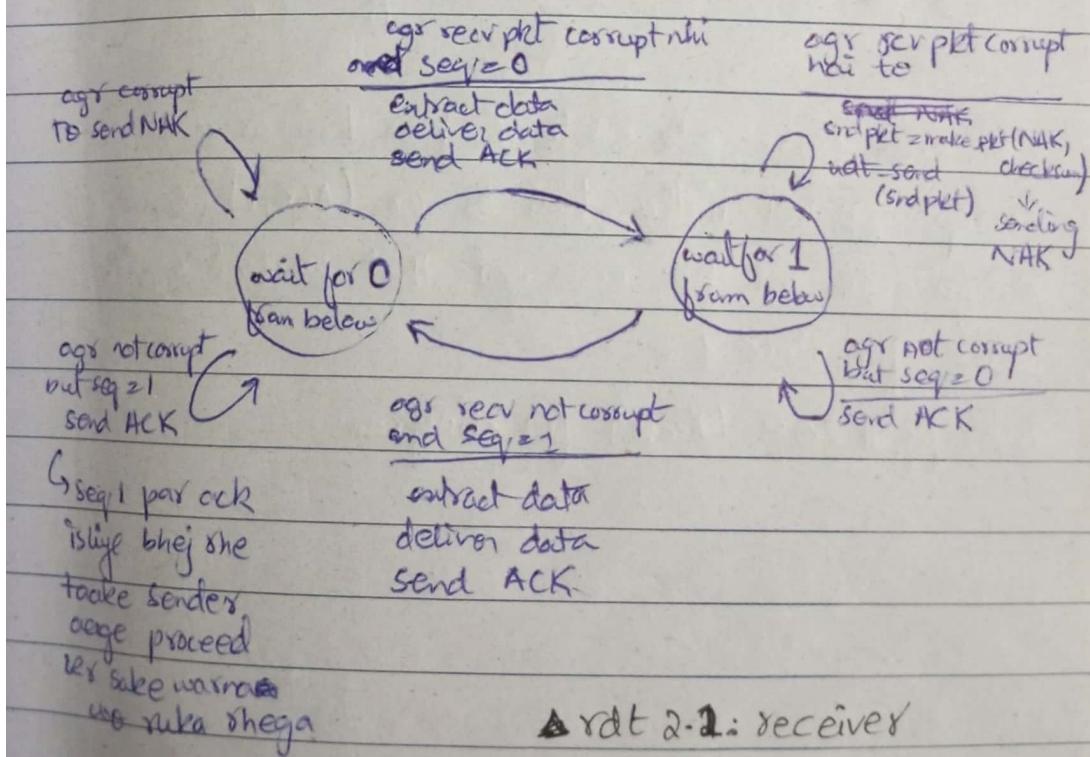
→ The sender knows that received NAK or ACK packet was generated in response to its most recently data packet.

Day

Date



RDT 2.1 : Sender



RDT 2.1 : receiver

Date

Day

Ydt 2.2: NAK free protocol

- instead of sending NAK, receiver sends ACK for last received correctly received packet.
- A sender that receives two ACKs for the same packet (means it receives duplicate ACKs) knows that the receiver did not correctly receive the packet following the packet that is being ACKed twice.

→ sender ne sequence no 0 ka packet send kia
hui to us (ACK,0) receive hua to hi aage
bhagega agr (ACK,1) receive ho means he
needs to retransmit. same for seq #1

→ receiver agr state 0 par hui or we
corrupt packet milta hui to us (ACK,1) send
krega so that he can get that packet again
same for state 1.

see Pg# 209 (book) for diagram.

- Date: _____ Day: _____
- alternating-bit protocol
- Qdt 3.0: Channels with error & Packet Loss
- ↳ Sender waits for a reasonable amount of time for ACK (time could be RTT)
 - ↳ if no ACK received retransmits
 - ↳ if packet or ACK is not lost but delayed
 - ↳ retransmission will be duplicate, but seq# already handle this
 - ↳ receiver must specify seq# of packet being ACKed.
- When sender sends packet it starts timer & when it receives ACK it stops it
- If sender receives an ACK, 1 when it is in state 0, it does not retransmit, instead it does nothing & only retransmits when it is timeout.
- Side pg # 64
- o Pipelined Reliable Data Transfer Protocol :
- RTT = 20 ms
 - R = 1 Gbps
 - L = 8000 bit
 - $d_{trans} = \frac{L}{R} = \frac{8000}{10^9} \text{ bit/sec} = 8 \mu\text{s}$
- It takes 8 microsec for sender to completely send its packet on channel.
- Page No.

Date _____

Day _____

→ Then he receives an ACK after 30ms so total ($t = RTT + L/R$) 30.008 msec time will be spent to sent just one. Then after 3.008ms the next packet will be sent.

→ Sender Utilization → fraction of time the sender is actually busy sending bits into channel.

$$U_{\text{sender}} = \frac{L/R}{RTT + L/R} = 0.00027$$

► Solution :-

↳ Rather than operate in a stop-and-wait manner, the sender is allowed to send multiple packets without waiting for ACK.

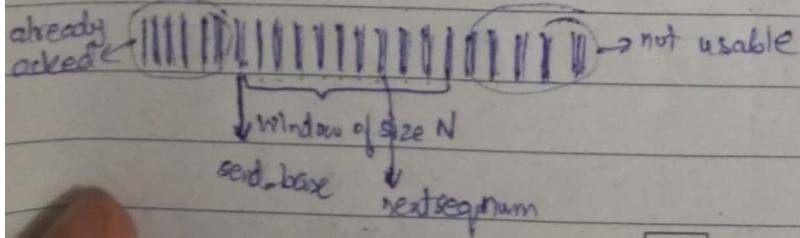
→ Agr ek waqt mein 3 packets soath bhej rhe hai to

$$U_{\text{sender}} = \frac{3L/R}{RTT + L/R} = \frac{0.0024}{30.008} = 0.00081$$

Go-Back-N :-

• Sender :-

is mein ek window size N hota hai that tells, consecutive transmitted but unAcked plots



Ack(n): Ack's all packets upto, including seq# n
 or ~~sakle~~<sup>ack receive hojaye 1st packet ki to window ek seq number aage move kr jayegi.
 → After timeout hojaye to nth packet or wke age jitne bhi window mein packet honge sakle & retransmit kr dega.</sup>

- receives & k end par agr ke packet nhi aata to jo last acknowledged packet hogega uski ack wapis send hojayegi
 out-of-order packets deal nhi honge.
- sender ke jab duplicate ACK receive hoji to wo wko ignore krdega or jo window mein jo packets honge unko transmit krdega.

Selective Repeat:

- Sender har packet ka timer maintain kr rha hota hai.
- agr ~~to~~ timeout hogaya to sfy wohi packet resend hoga.
- receiver ka par agr out of order receive ho to wo last acked plet ke baad wohi same packets ko buffer mein store krleta hai.
 like sender ne 0,1,2,3 send kia 0,1 pohch gaya or ack 1 ki rec hogai but 2 loss hogaya. Acha receive ka par 0,1 or 3 aya. 0 or 1 ka ack to wo send kr chuka tha lekin 2 nhi aya Page No. or 3 aayga to wene 3 ko buffer kr liya sender ka ja 2 ka timeout hogta to wapis resend or wohi ACK send kar di.

Date _____

Day _____

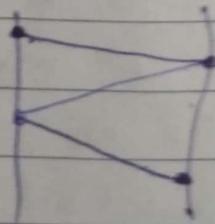
Kaha. Phir receiver ke job 2 receive kahe to aane 2
kei ACK bhejati or joct 2 & 3 ke application layer
ke deliver krdia.

→ Application layer ke kabhi bhi out-of-order packets
Send nahi kr skete.

"CONNECTION-ORIENTED TRANSPORT"

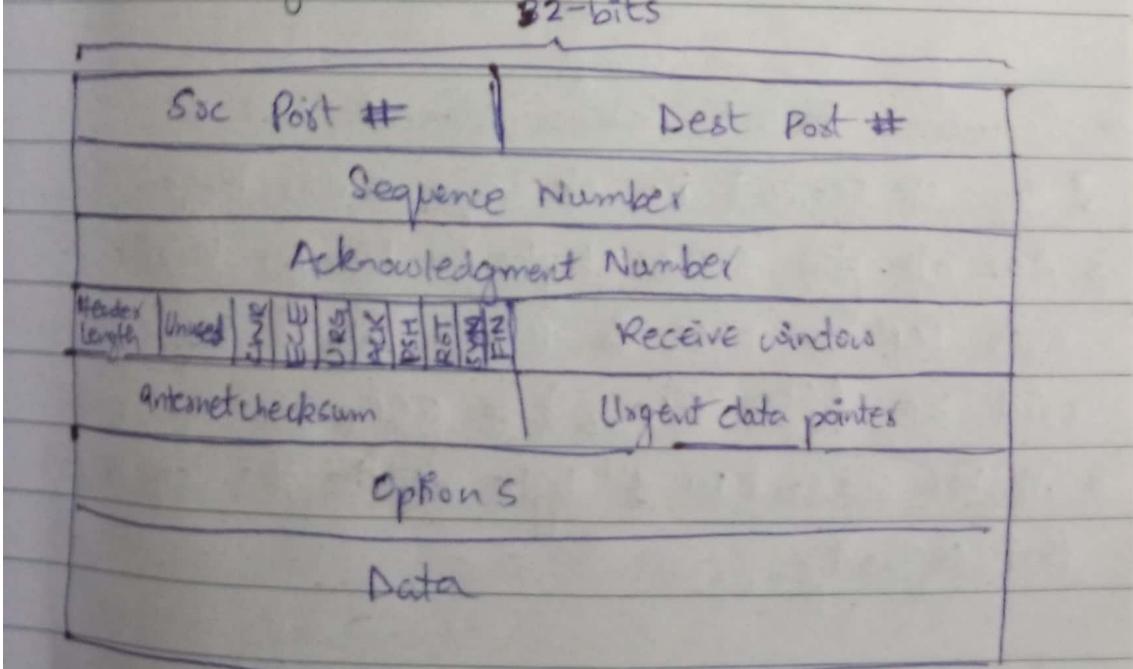
► TCP Connection :

- TCP connection provides a full duplex service.
↳ If there is a connection b/w A & B then
both can send & rec data.
- TCP connection is always point to point i.e b/w a
single sender & a single receiver.
↳
- Multicasting (transfer of data from one sender to
many receivers) is not possible in TCP.
- TCP Total 3 segments are sent sender rec
b/w sender/rec therefore this
connection establishment is called as
three way handshake



- Application layer se data TCP ko send buffer mein store hojata hai.
- ISI buffer se data network layer ko pass hata hai
- ER segment mein max katra data jai sakta hai,
 that is referred to as maximum segment size (MSS)
- Ye MSS hamare link layer ke frame ke size par depend karta hai that is called maximum transmission unit (MTU)
- Note: MSS is the max amount of application layer data in the segment, not the max size of TCP segment including headers.
- TCP pairs each chunk of client data with a TCP header, thereby forming TCP segment

► TCP Segment Structure



Date _____

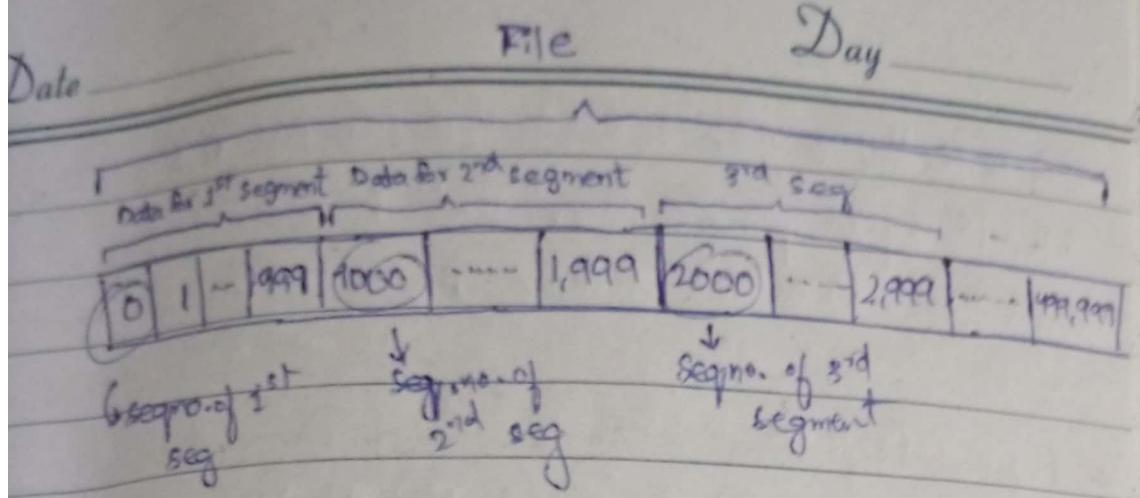
Day _____

Kisi phir receiver ka job 3 receive hoga
waise 2, 3 done ki ACK khej di.

- receive window (16-bit) used for flow control
- header length field (4-bit) specifies length of TCP header
- Options field
- Flag field (6 bits)
 - ↳ ACK used to indicate that the value carried in the ACK field is valid; i.e. the segment contains acknowledgement of a seq that was successfully received
 - ↳ RST,
 - ↳ SYN & FIN is used for connection setup & teardown
 - ↳ CWR & ECE used for congestion notification
 - ↳ PSH indicates the receiver should pass data to upper layer immediately
 - ↳ URG indicates that there is data in this seg that the sending layer has marked as "urgent"
 - ↳ the location of last byte of this urgent data is indicated by urgent data pointer field.

• Sequence No. & Acknowledgment No.s

- Har ek segment ka ek sequence no. hogा.
- Agr ek file 500,000 bytes ki transfer kرنा hai over TCP or MSS is 1000. Total segments will be $(500,000/1000) = 500$.
- Ab har segment ki 1st byte will be the seq.no of that segment.

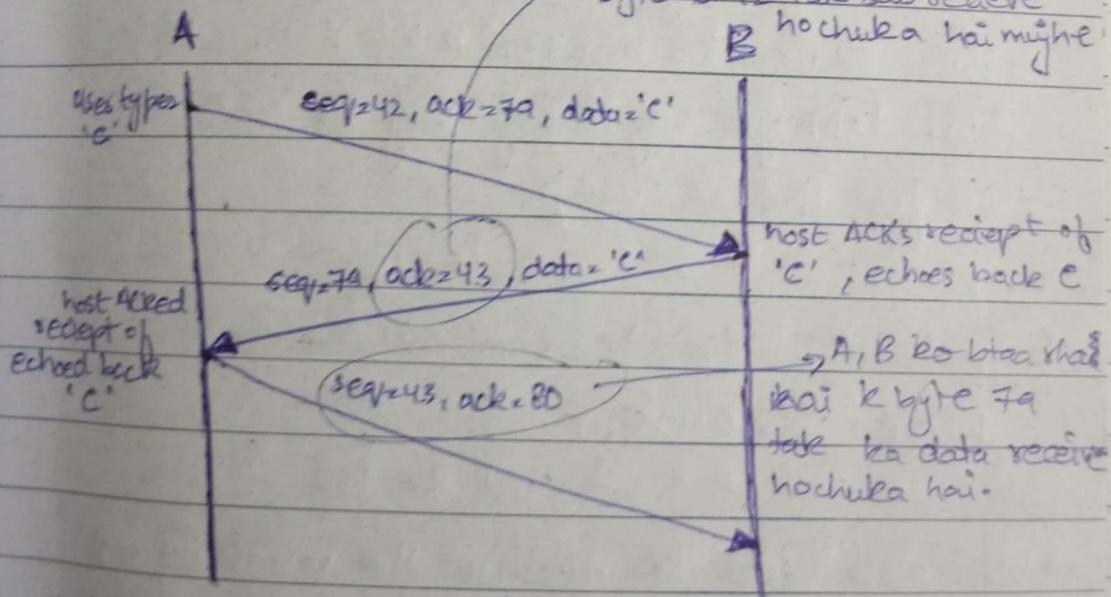


Acknowledgment Number:

The ackno. that Host A puts in its segment is the sequence no. of the next byte Host A is expecting from Host B.

Example: Agar A, B ko 0-999 bytes wala segment send kreg B ko to wo ack mein next segment ka seq no. i.e. 1000 send kregga.

→ B, A ko ye btaa rha hai k
byte no. 42 take sab receive
B hochuka hai mujhe



→ Client to server (jسے data) bhej rhe hain uska ack
server-to-client wale segment mein receive hota hai. This
is called piggybacked acknowledgment

Date _____

Day _____

→ Agr. out of order segments receive ho TCP mén
to its uplo programmer that how to deal
with it. i) either discard ii) or buffer

⇒ Round-Trip Time Estimation and Timeout :

- Timeout ki value hamesha RTT se basi honi
chahiye warna unnecessary retransmissions hongi
- Agr. timeout ki value baht zyada rakhdi te
agr. segment loss hogaya to retransmission
will be very slow.

- SampleRTT , for a segment is the amount of time
between when the segment is sent & when
an ack for segment is received.

Giske ek mastaka kisi seg ke liye measure ki jaata
hon. Lekin ye seg to seg vary kar sakte hai
so for that we use α

$$\text{EstimatedRTT} = (1-\alpha) * \text{Estimated RTT} + \alpha * \text{SampleRTT}$$

Where $\alpha = 0.125$

→ We also need to measure the variability of RTT

$$\text{DevRTT} = (1-\beta) * \text{DevRTT} + \beta * |\text{SampleRTT} - \text{Estimated RTT}|$$

Where $\beta = 0.25$

Date _____

Day _____

→ Now the final formula for timeout is

$$\text{Timeout Interval} = \text{Estimated RTT} + 4 * \text{Dev RTT}$$

safety margin

→ See slide # 19

- Doubling The Timeout Interval

→ Jab bhi timeout expire hota hai to new timeout
is twice the previous value (not calculated thru
formula)

→ Agar data application se receive hoa ho ya
acknowledgment receive hoofi ho tab hi sirf
formula se value derive hoti hai.

- Fast Retransmitt =

if sender receives 3 additional ACKs for same data
resend unmatched segment with smallest seq#

before
timeout

→ TCP Comparison with Go-Back N & Selective Repeat

↳ TCP's error recovery mechanism is a hybrid of
Go-Back-N & SR.

1. TCP is similar to GBN because

- It uses cumulative acknowledgments
- The sender tracks only the smallest unack seq#
- out of order seg are not immediately acked

2o TCP differs from GBN because

- Many TCP application implementations buffer out of order segments, unlike GBN
- It does not always retransmit all subsequent packets when a seg is lost.

↳ Example: Sender ne 1, 2, 3, ..., N segments send kro & the arrived correctly. But ACK for n^{th} segment is lost. Ako GBN hota to timeout hote hi $n^{\text{th}}, n+1, n+2, n+N$ packets retransmit hote. Lekin TCP sirf n^{th} packet transfer krega timeout pos. Jab tak n^{th} ki ACK receive nahi hogi $n+1$ seg ki retransmission nhi hoga. Jab ACK receive hogi to us mein boqi $n+1 - n+N$ sabki cumulative ACK aajayegi ~~kipha~~. And then it would not retransmit $n+1$ seg.

- Lekin agr timeout hone se phle $n+1$ ACK aajygi to n^{th} seg bhi retransmit nhi karna parega q k $n+1$ ACK ka mtlab hau k use phle k tamman seg recev ho chuke hau.

(n+1)

Sender ko agr ACK ~~mil~~ mil sha hau to iska matlab ye k k use phle tamman seg rec ho chuke hau. LEKN & GBN ACK(n+1) ko ignore kar deta hau.

3. TCP is similar to SR

- The receiver ACKs out-of-order segments individually.
- The sender retransmits only lost seg rather than unnecessary ones.

► Flow Control :

- TCP provides a "flow-control service" to its application to eliminate the possibility of the sender overflowing the receiver's buffer.
- It is like speed matching service
 - * matching the rate at which the sender is sending against the rate at which receiving application is ready"

→ The receive window ($rwnd$) field in TCP header

is used to control flow control.

→ it gives idea to sender that how much free space is available at receiver

$\boxed{\begin{array}{l} \text{\# of bytes} \\ \text{receiver willing to} \\ \text{accept} \end{array}}$

→ 4096 bytes

→ The receiver allocates a RcvBuffer to a connection, therefore in order to permit overflow of this buffer

$$\text{Last Byte Read}_{\text{Cvd}} - \text{Last Byte Read} \leq \text{RcvBuffer}$$

Now, the $rwnd$ is set to the amount of spare room in the buffer

$$rwnd = \text{RcvBuffer} - [\text{Last Byte Rcvd} - \text{Last Byte Read}]$$

Date _____

Day _____

- rwnd is dynamic variable as the size changes with each req.
- Initially the receiver sends $rwnd = \text{Recv Buffer}$
- The sender maintains,
$$\underbrace{\text{Last Byte Sent} - \text{Last Byte Acked}}_{\text{The amount of unACKed data that A has sent into the connection}} \leq rwnd$$

Scenario: Agar receiver ne $rwnd = 0$ send krdia
Sender ko or ab waise baad waise jaise koi data
nhi sender ko bhejne k liye (jisse new value of
 $rwnd$ is not provided to sender)
Ab sender koi data nhi bhejega qd $rwnd = 0$
hai. He is now blocked until the receiver sends
a non-zero value (but receiver half
nothing to send now).

In this case TCP requires sender to send segments
with one data byte when $rwnd = 0$. These seg
will be ACKed by receiver. Eventually buffer
will begin to empty & ACK will contain a
non-zero $rwnd$ value.

D TCP Connection Management :

• Initiating a Connection

① Client Sends SYN

i- The client sends a TCP SYN segment to the server to initiate a connection

ii- The SYN bit = 1, and the client randomly selects an initial seq# (client_isn)

iii- The segment is encapsulated in an IP datagram & sent to server

② Server Responds with SYN-ACK Segment

i- The server allocates buffers & variables for the connection

ii- Sends a SYN-ACK segment back to the client

iii- The SYN bit = 1, the ACK = client_isn + 1 & the server sends its own server_isn

③ Client Sends ACK

i- The client acknowledges the server's SYN-ACK by sending ACK segment

ii- The ACK = server_isn + 1, and SYN bit = 0 (as connection is established)

iii- This segment may carry application data.

Date _____

Day _____

- Now all future segments carrying app. data will have SYN bit = 0.
- This is called TCP Three-Way Handshake

Q: Why is a three-way handshake needed instead of a Two-way?

① Avoids Half Open Connection:

A two-way connection may lead to situations where one side thinks a connection is open while the other does not.

② Prevents Spoofing Attacks:

The 3rd step confirms that both parties received each other's seq#, reducing the risk of malicious connection hijacking.

③ A 2-way han would not confirm that server's response was received, leading to potential connection issues & packet loss.

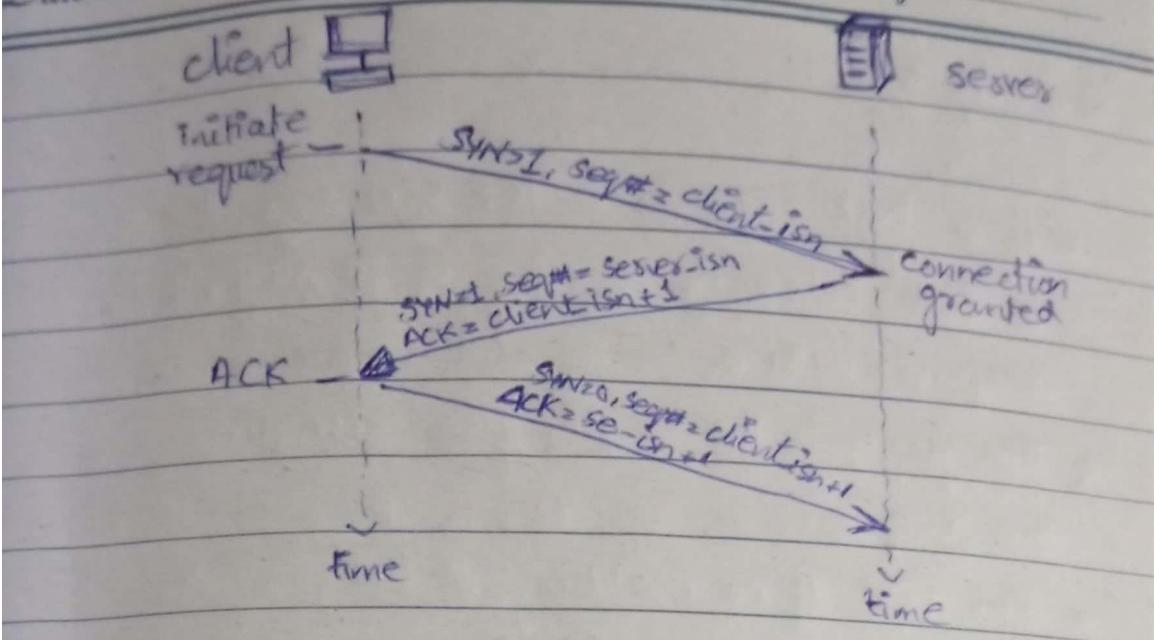
Q: Why are initial seq# needed?

① Unique initial seq# prevents prevents duplicate segments from prev connections from being mistaken as valid data.

② Ensures each TCP connection starts with unique seq#

Date _____

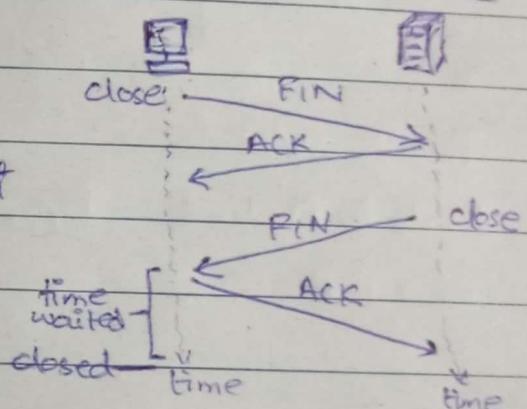
Day _____



• Closing a connection

① Client sends a FIN (Finish) segment to the server, indicating it wants to close the connection

② The server responds with an ACK segment confirming it rec the FIN request



③ Now the server sends a FIN seg to client, signaling that it is also ready to close conn.

④ Client sends a final ACK, config the termination

⑤ Both client & server de allocates resources,
connection terminated completely.

Date _____

Day _____

Scenarios :-

- If a host receives a TCP SYN segment for a port where no service is running, it responds with a RST (reset) segment
- This prevents the sender from resending the req.
- For UDP if no matching socket exist the host sends an ICMP "Port unreachable" message.

Day

Date

CHAPTER # 02 (Remaining Topics)

► PEER TO PEER FILE DISTRIBUTION

- The distribution time is the time it takes to get a copy of the file to all N peers.
- Peers req service from other peers, provide service in return to other peers.

Self-scalability — new peers bring new service capacity and new service demands

- File distribution Time : client-server
- server transmission : must sequentially send (upload) N file copies:
 - time to send one copy $\rightarrow \frac{F/U_s}{\text{server upload rate}}$
 - time to send N copies $\rightarrow NF/U_s$
- each client must download file copy
 $d_{min} = \text{min client download rate}$
 $\text{min client download time} \rightarrow F/d_{min}$

Time to distribute

$$F \text{ bits to } N \text{ clients } D_{cs} \geq \max \left\{ \frac{NF}{U_s}, \frac{F}{d_{min}} \right\}$$

using client-server

approach

Increases linearly with
no of peers/clients

- File distribution time : P2P

- Server transmission : must upload at least one copy

$$\text{time to send one copy} \rightarrow \frac{F}{U_s}$$

- client : each client must download file copy
min client download time $\rightarrow \frac{F}{d_{\min}}$

- clients : The system must deliver (upload) F bits to each of the N peers, thus delivering total of NF bits so,

$$\text{max upload rate} \rightarrow \frac{NF}{U_s + \sum U_i}$$

Time to distribute

F bits to N clients

using P2P

$$D_{P2P} \geq \max \left(\frac{F}{U_s}, \frac{F}{d_{\min}}, \frac{NF}{U_s + \sum U_i} \right)$$

ye bhi linearly
grow krega reka

ye bhi grow krega
to who nullify kroga
reduce
9/12 wo uploading
service bhi saath
layega.

See slide #86 for BitTorrent File Distribution

▷ VIDEO STREAMING :

- Video is stored on an HTTP server as a regular file with a URL
- The client sends an HTTP GET request to fetch the video
- The server responds by sending the video as fast as network conditions allow
- The client buffers the received bytes and starts playback once a threshold is reached
- The streaming application decompresses and displays video frames as they arrive

• Limitations of Basic HTTP Streaming

- Fixed encoding: All clients rec same video quality regardless of bandwidth difference
- Bandwidth variation issues : If a user's connection is slow, buffering & playback interruptions occur

• DASH (Dynamic Adaptive Streaming over HTTP)

- Video is encoded into multiple versions at different bitrates & quality levels
- Client request video in chunks of a few seconds instead of whole file
- Clients dynamically chooses higher or lower bitrates based on available bandwidth.

Date _____

Day _____

- Users with high speed internet get better quality
- Implementation
 - Each version of the video is stored on the HTTP server with a different URL
 - The server provides a manifest file containing URLs & bitrates of all versions
 - The client 1st downloads the manifest file, then selects chunks based on avail. bandw.
 - Client measures received bandwidth & run a rate determination algo to decide the next chunk.

② Content Distribution Network (CDNs)

- Challenges in centralized Video Streaming
 - If data center goes down, streaming stops entirely.
(Single point of failure)
 - More links incr the risk of freezing delays due to bandwidth constraints
 - Popular videos are repeatedly sent over the same network, increasing cost

• Role of CDN

- CDN stores copies of videos in multiple geographically

distributed servers.

→ Reduces latency by serving users from the closest available servers

→ CDNs can be private

- CDN Server Placement Strategies

- Enter Deep

- ↳ Deploys small server clusters inside ISPs worldwide to minimize distance to users

- ↳ Improves performance but is complex to maintain due to thousands of locations

- Bring Home

- ↳ Deploys fewer, larger clusters at Internet Exchange Points (IXPs) instead of inside ISPs

- ↳ Easier to manage but may increase delay & lower throughput

- Content Replication & Caching in DNS

- Not all videos are stored in every cluster to optimize storage.

- Uses a pull strategy: A requested video not in a cluster is fetched from another cluster or a central repository

- Least frequent used videos are removed when storage is full.

Date _____

Day _____

- Cluster Selection Strategies

- 1. Geographical Proximity based Selection

- ↳ Uses geo-location databases to find the closest server by distance

- ↳ Drawbacks:

- ↳ The nearest server geographically may not be the best in terms of network latency or hops

- ↳ Some users use remote DNS servers, leading to incorrect mapping

- 2. Real Time Network Performance-Based Selection

- ↳ CDNs periodically measure delay & packet loss b/w clusters & local DNS server

- ↳ Uses ping, DNS queries to dynamically choose best cluster

- ↳ Drawbacks:

- ↳ Many LDNS servers block probes, making it hard to collect accurate data.

a small test message
sent by CDN to
measure network per-
for

Date _____

Day _____

CASE STUDY: Netflix

- Uses Amazon Cloud for web services
- Uses its own private CDN for video distribution

- Amazon Cloud Functions

- Content Ingestion : Uploads studio master versions of movies
- Content Processing : Converts movies into multiple formats for different devices & bitrates
- Uploading to CDN : Processed versions are sent to Netflix's CDN

- Netflix's Private CDN

- Servers are installed in IXPs & ISPs
- Over 200 IXP locations & hundreds of ISPs host Netflix servers
- Server contains entire or most popular videos
- Push caching & content is proactively pushed to servers during off peak hrs.

- Client - Server Interaction

- Web browsing is handled by Amazon Cloud Services
- When a user selects a movie:
 - The best CDN server is chosen based on loc & availability.

- If a user's ISP has a Netflix CDN back, it is preferred, otherwise a nearby ISP CDN is used
 - The client recs the server's IP addrs. and a manifest file
 - The client interacts directly with the CDN server using a proprietary version of DASH
 - Uses HTTP GET byte-range request to fetch 4-second video chunk
 - Client measures throughput and adjust video quality dynamically
-
- Netflix CDN design Advantage
 - The client is directly assigned a CDN server instead of resolving via DNS
 - Unlike general purpose CDN, it is dedicated to video delivery, simplifying its design.

CASE STUDY: YOUTUBE

- Uses its own private CDN
- Pull caching is used, meaning videos are fetched on demand & stored temporarily.
- DNS redirection helps balance traffic by assigning users to optimal server based on RTT
- uses HTTP streaming but does not support DASH