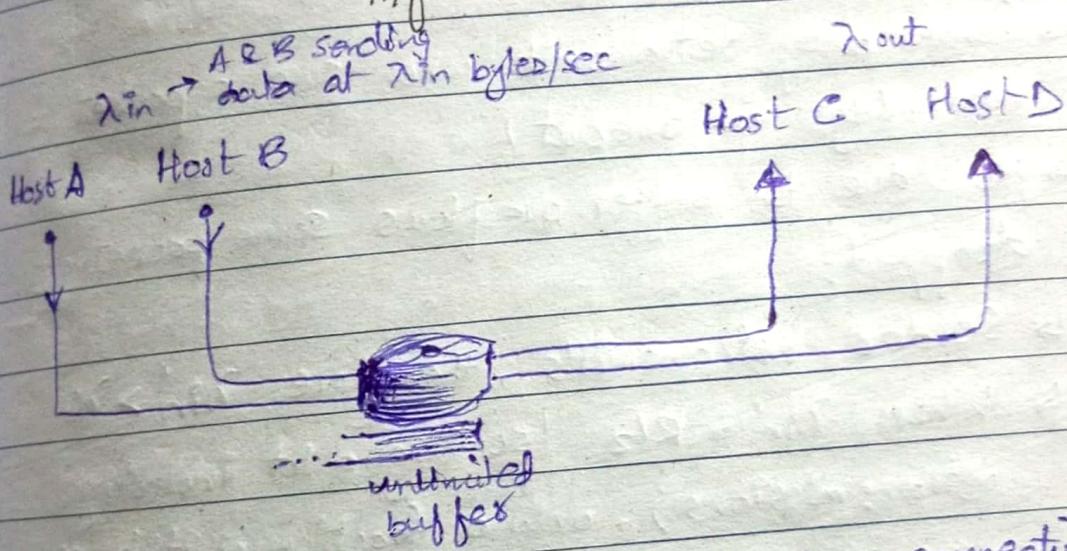


► PRINCIPLES OF CONGESTION CONTROL

- "Too many sources sending too much data too fast for network to handle" — congestion
- Symptoms of congestion control
 - ↳ long delays ↳ packet loss

• Scenario #01: Two Senders, a Router with Infinite Buffer

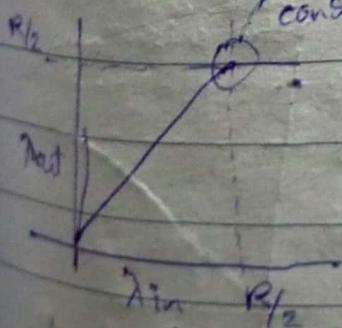


Agr link capacity R hui to ~~per connection~~ per connection
 $R/2$ hogi:

→ Maximum sender ki sending capacity $R/2$ hui

hosakti hui.

is point $R/2$
constant

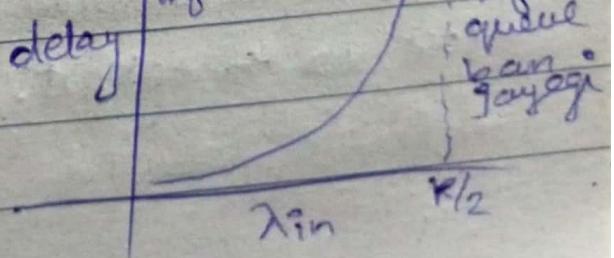


Page No. []

agr hui to $R/2$ par send karne lage to delay aa hujega a.

ek buffer infinite hui

is ek infinite queue kar gayegi



Date _____

Day _____

- Scenarios #02: Two Senders and a Router with Finite Buffers

→ Case #01:

↳ Sender knows when buffer is free and only sends then.

↳ No packet loss

↳ So throughput = $\lambda_{in}' = \lambda_{in}$

↳ Max throughput/host = $R/2$

original data
+ retransmission

→ Case #02:

↳ Sender only retransmits when sure a packet is lost (e.g. via timeout)

↳ Now due to buffer overflows some packets are dropped.

↳ Sender retransmits, therefore $\lambda_{in}' > \lambda_{in}$

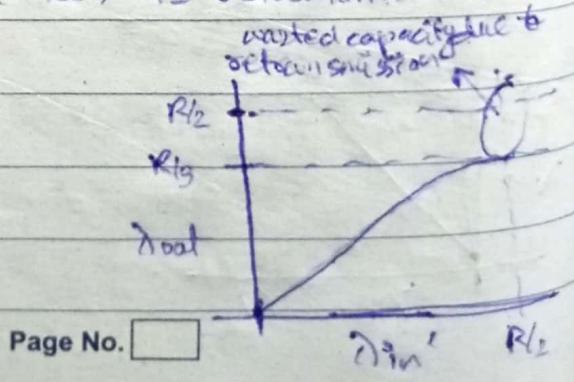
↳ When $\lambda_{in}' = R/2$ (sending a lot including retransmissions), only a fraction of it is new original data

↳ Delivered throughput = $R/3$

↳ $R/2$ means $R/3$ isn't original data but

↳ baagi $R/6$ (0.166) is retransmitted

↳ So, retransmission cost us throughput.

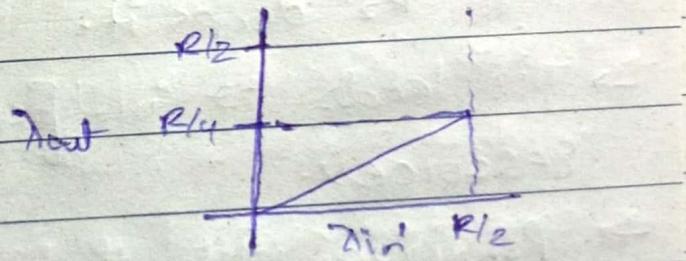


Date

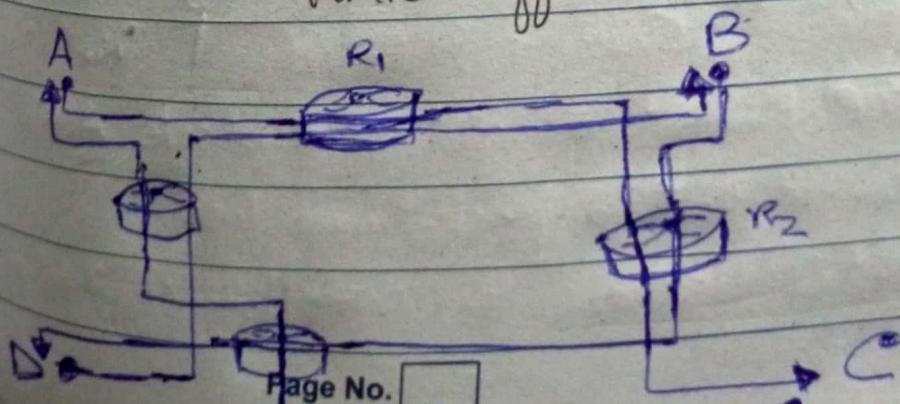
→ "sender must perform retransmission in order to compensate for dropped packets due to buffer overflow."

→ Case #03 =

- ↳ Sender mistakenly thinks a packet is lost due to delay & retransmits it.
- ↳ But in reality original pkt is still in queue.
- ↳ The router wastes link bandwidth forwarding duplicate pkt.
- ↳ This bandwidth could have been used to send useful pkt.
- ↳ When each pkt is forwarded twice on avg: throughput becomes $R/4$ when load is $R/2$



• Scenario #03: Four senders and a Router with Finite Buffer & Multi-hop Paths



Date _____

Day _____

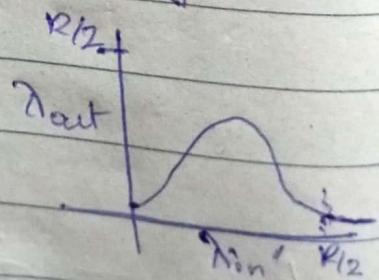
→ High Traffic :

- Consider A → C path & passes thru $R_1 \rightarrow R_2$
- Also, D → B shares R_1 , and B → D shares
- As λ_{in} increases for all.

R_2 ka buffer B-D connection se phale hi

jill hogaega apne R_1 ka 1st router R_2 ka
tablee A ke liye R_2 2nd router hua.

- A-C pkts arriving from R_1 are frequently dropped at R_2 .
- Eventually AC throughput drops to 0, despite sending more data.
- Packets dropped at 2nd hop router waste all upstream efforts.
 - ↳ work done at R_1 to send AC packets is wasted if R_2 drops it.
- This leads to inefficient use of network resources.
- Polarize pkts that have already traversed hops.



• Costs of Congestion

- ① Throughput can never exceed capacity.
- ② Delay increases as capacity approached
- ③ Loss / retransmission decreases effective throughput
- ④ Un-needed duplicates further decreases effective throughput
- ⑤ Upstream transmission capacity / buffering wasted for pkt lost downstream.

• Approaches Towards Congestion Control

① End-To-End Congestion Control

↳ Network layer congestion k baare mein kuch nahi btati

↳ TCP ko khud infer kina parega k network mein congestion hua ya nahi. By the help of packet loss & delay.

② Network assisted Congestion Control

↳ Routers (Network Layer) provide explicit feedback to sender & rec about congestion.

↳ May indicated congestion level or explicitly set sending rate

• TCP CONGESTION CONTROL :

- TCP controls the sending rate based on network congestion.
- If path is clear, it sends faster; if it's congested, it slows down.

Q: How TCP limits sending Rate

- TCP uses a variable called the congestion window (cwnd)
- The sender can't have more than "cwnd" bytes of unack data in the network.

$$\text{LastByte Sent} - \text{LastByte Acked} \leq \min(\text{cwnd}, \text{rwnd})$$

↑
ignore this

- TCP sends "cwnd" bytes per RTT
- send rate \leq cwnd / RTT
- Cwnd directly affects how fast TCP can send data
- congestion control is about dynamically adjusting cwnd based on perceived network conditions

Q: How TCP detects congestion?

1. Loss Event = Congested Signal

- If TCP times out or get 3 duplicate ACKs, it assumes that a packet was lost
- Packet is lost due to overflow of routes buff

- Buffer is overflowed when there is network congestion.
- On packet loss \rightarrow Decrease Rate
 - \hookrightarrow Loss means congestion
 - \hookrightarrow TCP reduces cwnd to reduce sending rate

- a. ACKs = Network is healthy
- If ACKs are arriving steadily.
 - TCP assumes no congestion & gradually incr. its send rate
 - TCP increases cwnd, probing for how much bandwidth it can safely use.

\rightarrow Self-Clocking Mechanism

- TCP uses ACKs to trigger data sending
- More/Faster ACKs = more data sent
- This automatically adjust the rate based on network feedback.

\rightarrow Bandwidth Probing Strategy

- TCP keeps increasing rate (cwnd) until pkt loss.
- Then it backs off, then slowly incr. again

► TCP Congestion Control Algorithm

The algorithm has 3 major component

- ① Slow Start
 - ② Congestion Avoidance
 - ③ Fast recovery
- mandatory components of TCP,
differing in how they incr
size of cwnd in response
to ACKs.

① Slow Start

→ Slow start alg increases the sending rate exponentially.

a. Initial Setup

- cwnd starts at 1 MSS (Max seg size)
- Only 1 segment is sent initially.

b. ACKs Double cwnd

- After each ACK receive none per 1 seg increase in cwnd.
- ~~This~~ Jiski waja se exponential growth happens
- $1 \rightarrow 2 \rightarrow 4 \rightarrow 8 \rightarrow \dots$ per RTT.

c. When does slow start end?

- If timeout

↳ set cwnd = 1

slow start threshold

↳ set ssthresh = cwnd/2

↳ Start slowstart again

- If $cwnd > ssthresh$

↳ Switch to Congestion Avoidance (grows linearly)

- If 3 duplicate ACK's
 - ↳ Enter Fast Recovery instead of restarting

② Congestion Avoidance

- ↳ It starts when the slow start ends typically when cwnd \geq ssthresh
- ↳ It takes more cautious, linear approach to avoid congestion
- ↳ How cwnd grows?
 - cwnd increases 1 MSS per RTT
 - cwnd increases by $MSS/cwnd$ per ACK
 - ↳ $MSS = 1,460 \text{ bytes}$, $cwnd = 14,600 \text{ bytes}$
 - ↳ No. seg = 10 ($MSS/cwnd$)
 - ↳ Each ACK increase cwnd by $1/10 \text{ MSS}$
 - When congestion is detected
 - Timeout:
 - ↳ Set cwnd = 1 & ssthresh = cwnd/2
 - ↳ Start slow start again
 - 3 Dup ACKS:
 - ↳ ssthresh = cwnd/2
 - ↳ cwnd = ssthresh + 3 * MSS
 - ↳ Enter Fast Recovery

→ 3 dup ACK fast recovery is like loss of 3 dup ACKs
 ↳ Matlab kuch segments rec hogai hai sirf ek whi hog. Matlab network completely congested nahi.
 To aggressively

③ Fast Recovery

↳ It is triggered on 3 dup ACKs.

↳ Action :

- Retransmit the lost segment immediately.
- Set ssthresh = cwnd/2 ^{plus 3 dup segments}
- set cwnd = ssthresh + 3 * MSS
- For each additional duplicate ACK, increase cwnd by 1 MSS -
3. dup ACK k ilawa ays ACK aye.

→ Exit Condition

↳ When ACK for lost seg arrives → enter congestion avoidance and reset cwnd = ssthresh

→ When timeout → go to slowstart ($cwnd=1$)
MSS

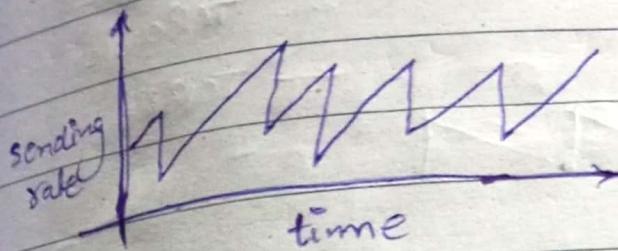
→ Earlier version of TCP, known as TCP Tahoe unconditionally sets $cwnd = 1$ MSS on timeout or 3 dup ACKS.

→ The newer version TCP Reno incorporated fast recovery.

Day

Date

- TCP congestion control is often referred to as an additive-increase, multiplicative-decrease!
- incr sending rate by 1 MSS / RTT until loss detected
- cut sending rate in half at each loss event



• TCP Cubic &

- TCP Reno's additive-increase & multiplicative decrease is too cautious after first loss
- Slow growth after cutting the cwnd may underutilize bandwidth
- CUBIC fixes this by faster recovering to the pre-loss sending rate & then probing carefully

→ How it works?

h_{max} = cwnd size when last loss occurred

K = Time needed (from loss) to reach h_{max} again

t = current time

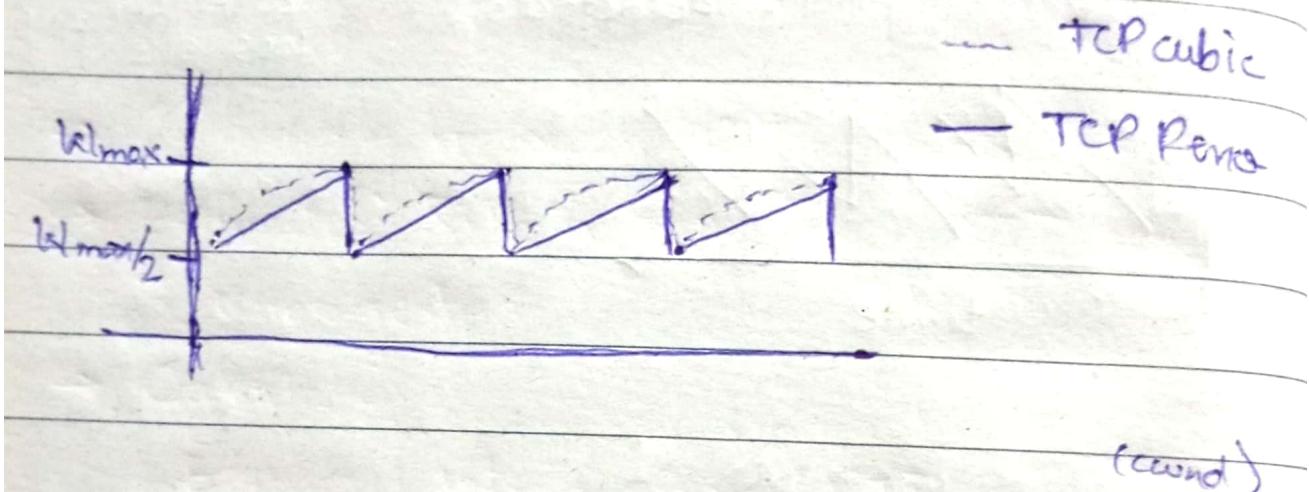
Increases h_t as a function of the cube of the distance between current time & K

Date _____

Day _____

↳ larger increases when further away from
K

↳ smaller increases (cautious) when near to
K



→ avg throughput of a = $\frac{0.75 \cdot W}{RT}$?
connection

► NETWORK ASSISTED CONGESTION CONTROL

① Explicit Congestion Notification

→ It allows routers to signal congestion without dropping pkts.

→ How ECN Works?

↳ IP Layer (Network layer)

- Uses 2 bits in the IP Header's Type of service (ToS) field.

- Router marks pkts instead of dropping them if congestion is detected

- Marked pkts are forwarded to receiver.

↳ TCP Layer (Transport layer)

- Receiver detects ECN marking and sets the ECE (ECN Echo) bit in the ACK

- Sender, upon rec ECE, reduces congestion window to half.

- Sender then sets the CWR (Congestion Window Reduced) bit in the next outgoing segment.

→ Proactive cong. control = Acts before pkt loss occurs

↳ Reduces retransmission

② Delay Based Congestion Control

→ Proactively detects congestion by monitoring increasing delay (RTT) rather than waiting for pkt loss

• TCP Vegas Approach

→ Continuously measures RTT of packets

→ Estimates RTT_{min} = minimum RTT when path is uncongested

→ Calculates:

$$\text{Expected rate} = cwnd / RTT_{min}$$

Actual rate = measured from ACKs # bytes sent in ACK / RTT measured

→ if Actual < Expected → congestion detected

→ Decrease cwnd

→ if Actual \geq Expected → Network uncongested

→ Increase cwnd

▷ TCP Fairness

→ If K TCP connections share a bottleneck link of rate R, fairness means each connection should ideally get R/K bandwidth.

→ Assumptions:

↳ Same MSS and RTT for both connections

↳ Operating in Cong. Avoidance mode

↳ Both conn use AIMD

- throughput Behaviour
 - ↳ Each TCP connection incr cwnd by 1 MSS / RTT
 - ↳ throughput of both increases together until loss occur
 - ↳ On pckt loss, both halve their cwnd
 - ↳ throughput drop
 - ↳ This cycles repeat, causing throughput to fluctuate near fairness line
 - ↳ Both will share equal bandwidth

- Fairness can be violated if
 - ↳ Connections have diff RTTs
 - ↳ Shorter RTTs = faster incr = more bandwidth
 - ↳ Presence of UDP traffic which doesn't reduce rate on congestion.

▷ Fairness in UDP

- Multimedia apps use UDP bcz they don't want rate throttled by congestion control
- Send pkts at constant rate
- Can tolerate pckt loss
- It is possible for UDP sources to crowd out TCP traffic